

Lecture 14

Recognition and kNN

Administrative

A3 is out

- Due Feb 21st

A4 is out

- Due Mar 7th

Administrative

Recitation this friday

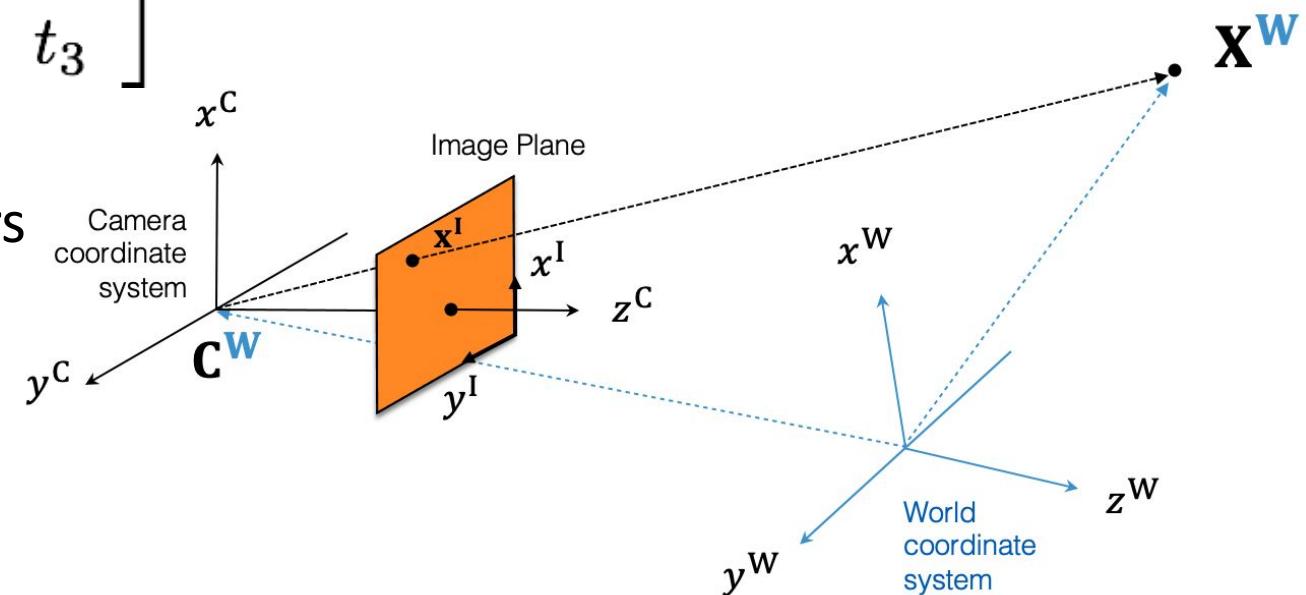
- Ontologies

So far: General pinhole camera matrix

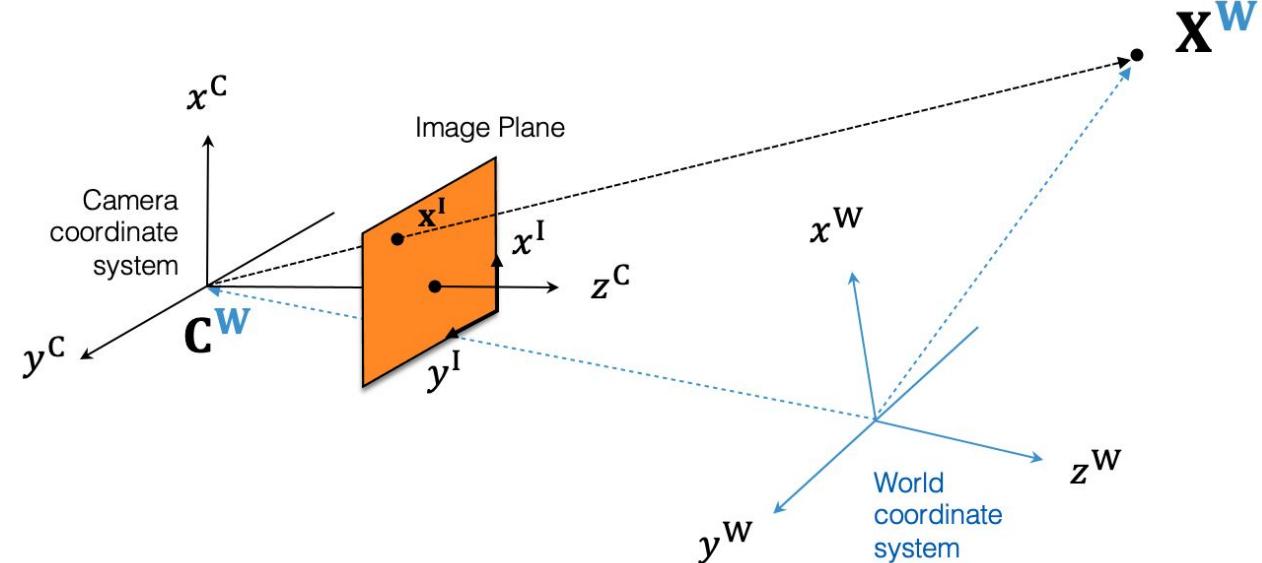
$$\mathbf{P} = \mathbf{K}[\mathbf{R}|\mathbf{t}] \quad \text{where} \quad \mathbf{t} = -\mathbf{R}\mathbf{C}$$

$$\mathbf{P} = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_1 & r_2 & r_3 & t_1 \\ r_4 & r_5 & r_6 & t_2 \\ r_7 & r_8 & r_9 & t_3 \end{bmatrix}$$

intrinsic parameters extrinsic parameters



So far: The Pinhole Camera Model



$$\tilde{\mathbf{x}}^I \sim \mathbf{P} \tilde{\mathbf{x}}^W$$

$$\mathbf{P} = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix} [\mathbf{I} \quad | \quad \mathbf{0}] \begin{bmatrix} \mathbf{R} & -\mathbf{R}\mathbf{C} \\ \mathbf{0} & 1 \end{bmatrix} = \mathbf{K}[\mathbf{R}|\mathbf{t}]$$

intrinsic parameters \mathbf{K} (3×3):
correspond to camera
internals (image-to-image
transformation)

perspective projection (3×4):
maps 3D to 2D points
(camera-to-image
transformation)

extrinsic parameters (4×4):
correspond to camera externals
(world-to-camera
transformation)

So far: Solving for camera matrix via total least squares

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x}} \|\mathbf{A}\mathbf{x}\|^2 \text{ subject to } \|\mathbf{x}\|^2 = 1$$

$$\mathbf{A} = \begin{bmatrix} \mathbf{X}_1^T & 0 & -\mathbf{x}'_1 \mathbf{X}_1^T \\ \mathbf{0} & \mathbf{X}_1^T & -\mathbf{y}'_1 \mathbf{X}_1^T \\ \vdots & \vdots & \vdots \\ \mathbf{X}_N^T & 0 & -\mathbf{x}'_N \mathbf{X}_N^T \\ \mathbf{0} & \mathbf{X}_N^T & -\mathbf{y}'_N \mathbf{X}_N^T \end{bmatrix} \quad \mathbf{x} = \begin{bmatrix} p_1 \\ p_2 \\ p_3 \end{bmatrix}$$

$$\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^\top$$

Equivalently, solution \mathbf{x} is the Eigenvector corresponding to the smallest Eigenvalue of \mathbf{A}

So far: Decomposition of the Camera Matrix

$$\mathbf{P} = \left[\begin{array}{c|c} & \bar{\mathbf{P}} \\ \hline & \begin{array}{c} p_4 \\ p_8 \\ p_{12} \end{array} \end{array} \right] \sim \mathbf{K} [\mathbf{R} | - \mathbf{R}\mathbf{C}]$$

$$\bar{\mathbf{P}}^T \bar{\mathbf{P}} \sim \mathbf{K}^T \mathbf{K} \text{ with } \mathbf{K} \text{ upper triangular p.d.}$$

Obtain \mathbf{K} by [Cholesky decomposition](#) of $\bar{\mathbf{P}}^T \bar{\mathbf{P}} = \mathbf{L} \mathbf{L}^T$

$$\mathbf{K} \sim \mathbf{L}^T$$

$$|\mathbf{R}| = 1 \Rightarrow \lambda = |\mathbf{K}^{-1} \bar{\mathbf{P}}|^{-1/3}$$

Once \mathbf{K} is known, we can compute $\mathbf{R} = \lambda \mathbf{K}^{-1} \bar{\mathbf{P}}$

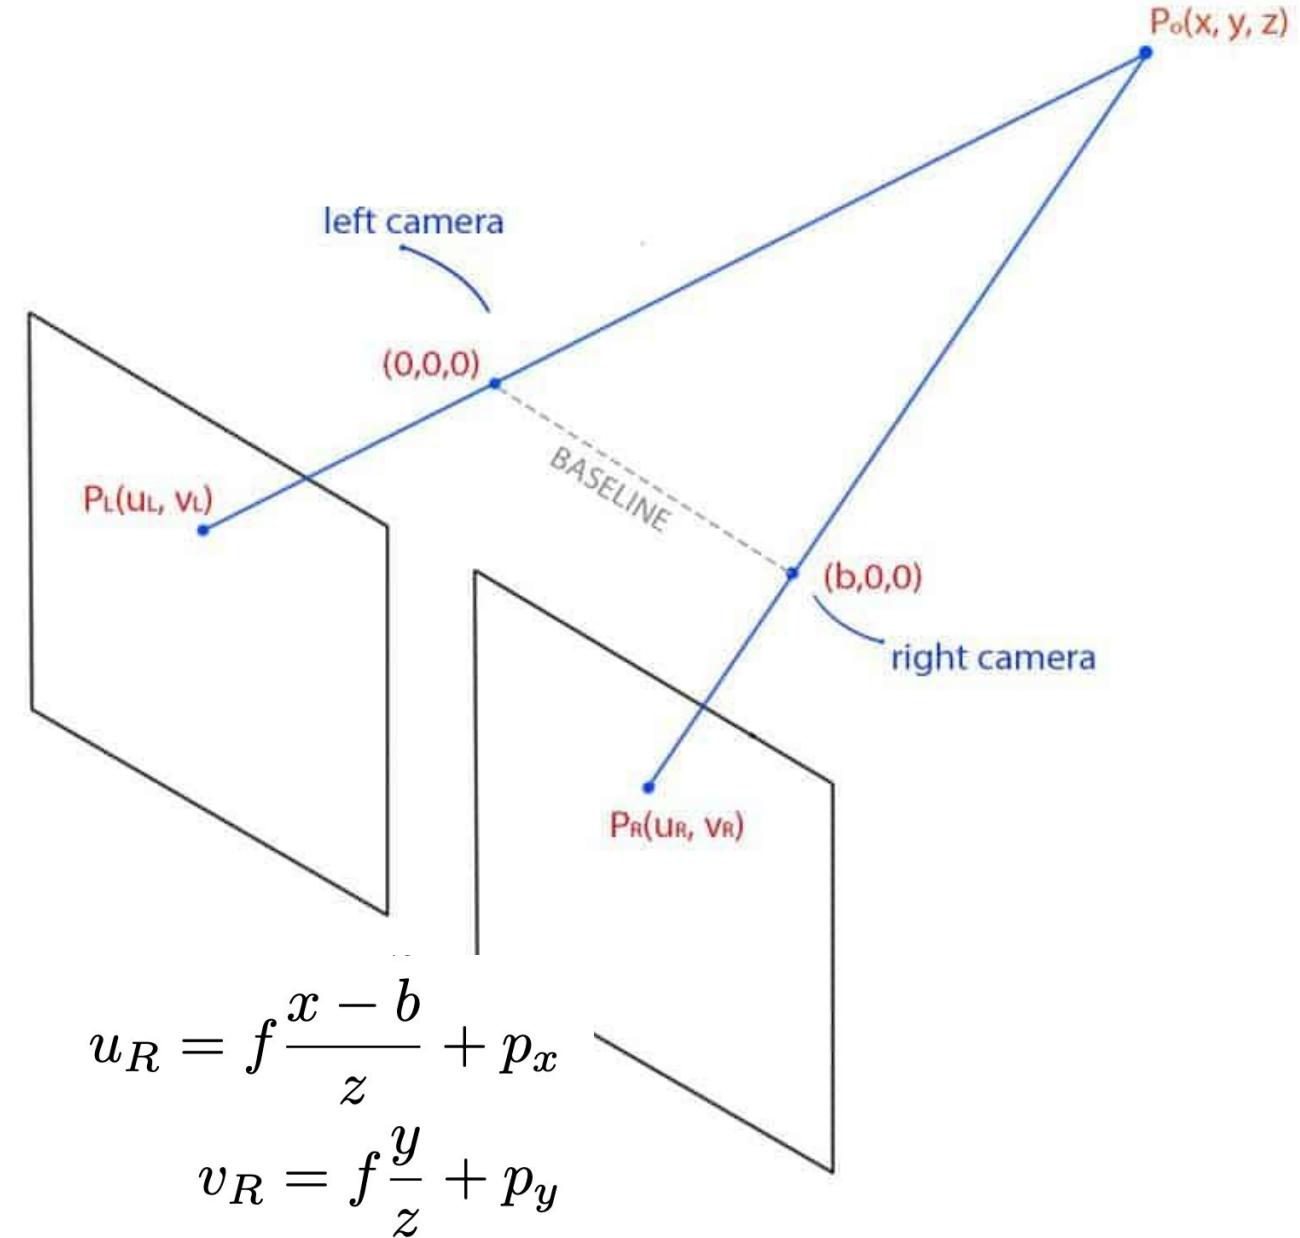
Finally, easy to know the camera center: $\mathbf{C} = -\lambda^{-1} \mathbf{R}^T \mathbf{K}^{-1} [p_4 \ p_8 \ p_{12}]^T$

So far: Estimating depth

$$u_L = f \frac{x}{z} + p_x$$

$$v_L = f \frac{y}{z} + p_y$$

$$K = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix}$$



Today's agenda

- Introduction to recognition
- A object recognition pipeline
- Choosing the right features
- A training algorithm: KNN
- Testing an algorithm
- Challenges with kNN
- Dimensionality reduction

Today's agenda

- Introduction to recognition
- A object recognition pipeline
- Choosing the right features
- A training algorithm: KNN
- Testing an algorithm
- Challenges with kNN
- Dimensionality reduction

What do we mean by recognition?



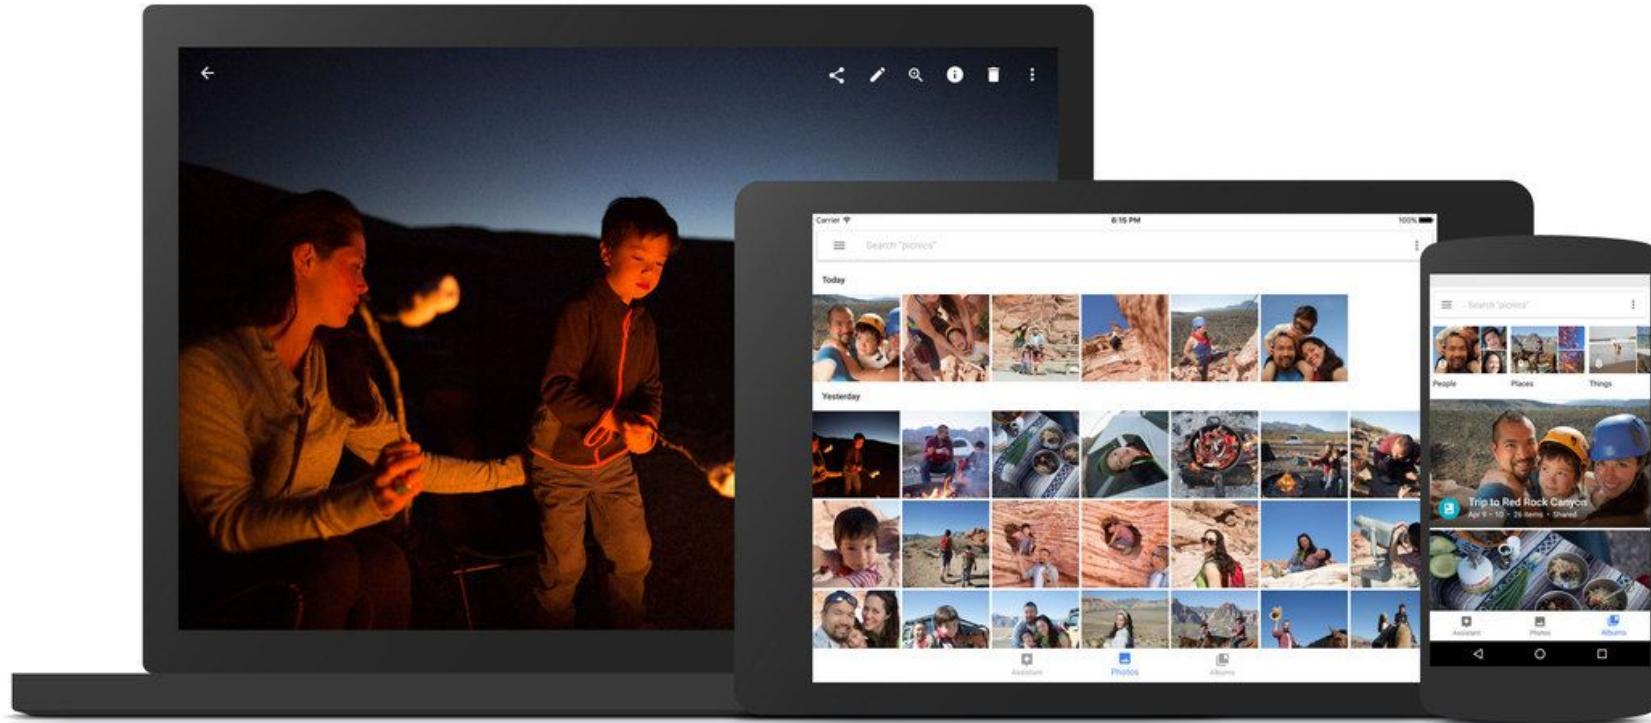
Classification: Does this image contain a building? [yes/no]



Classification: Is this an beach?



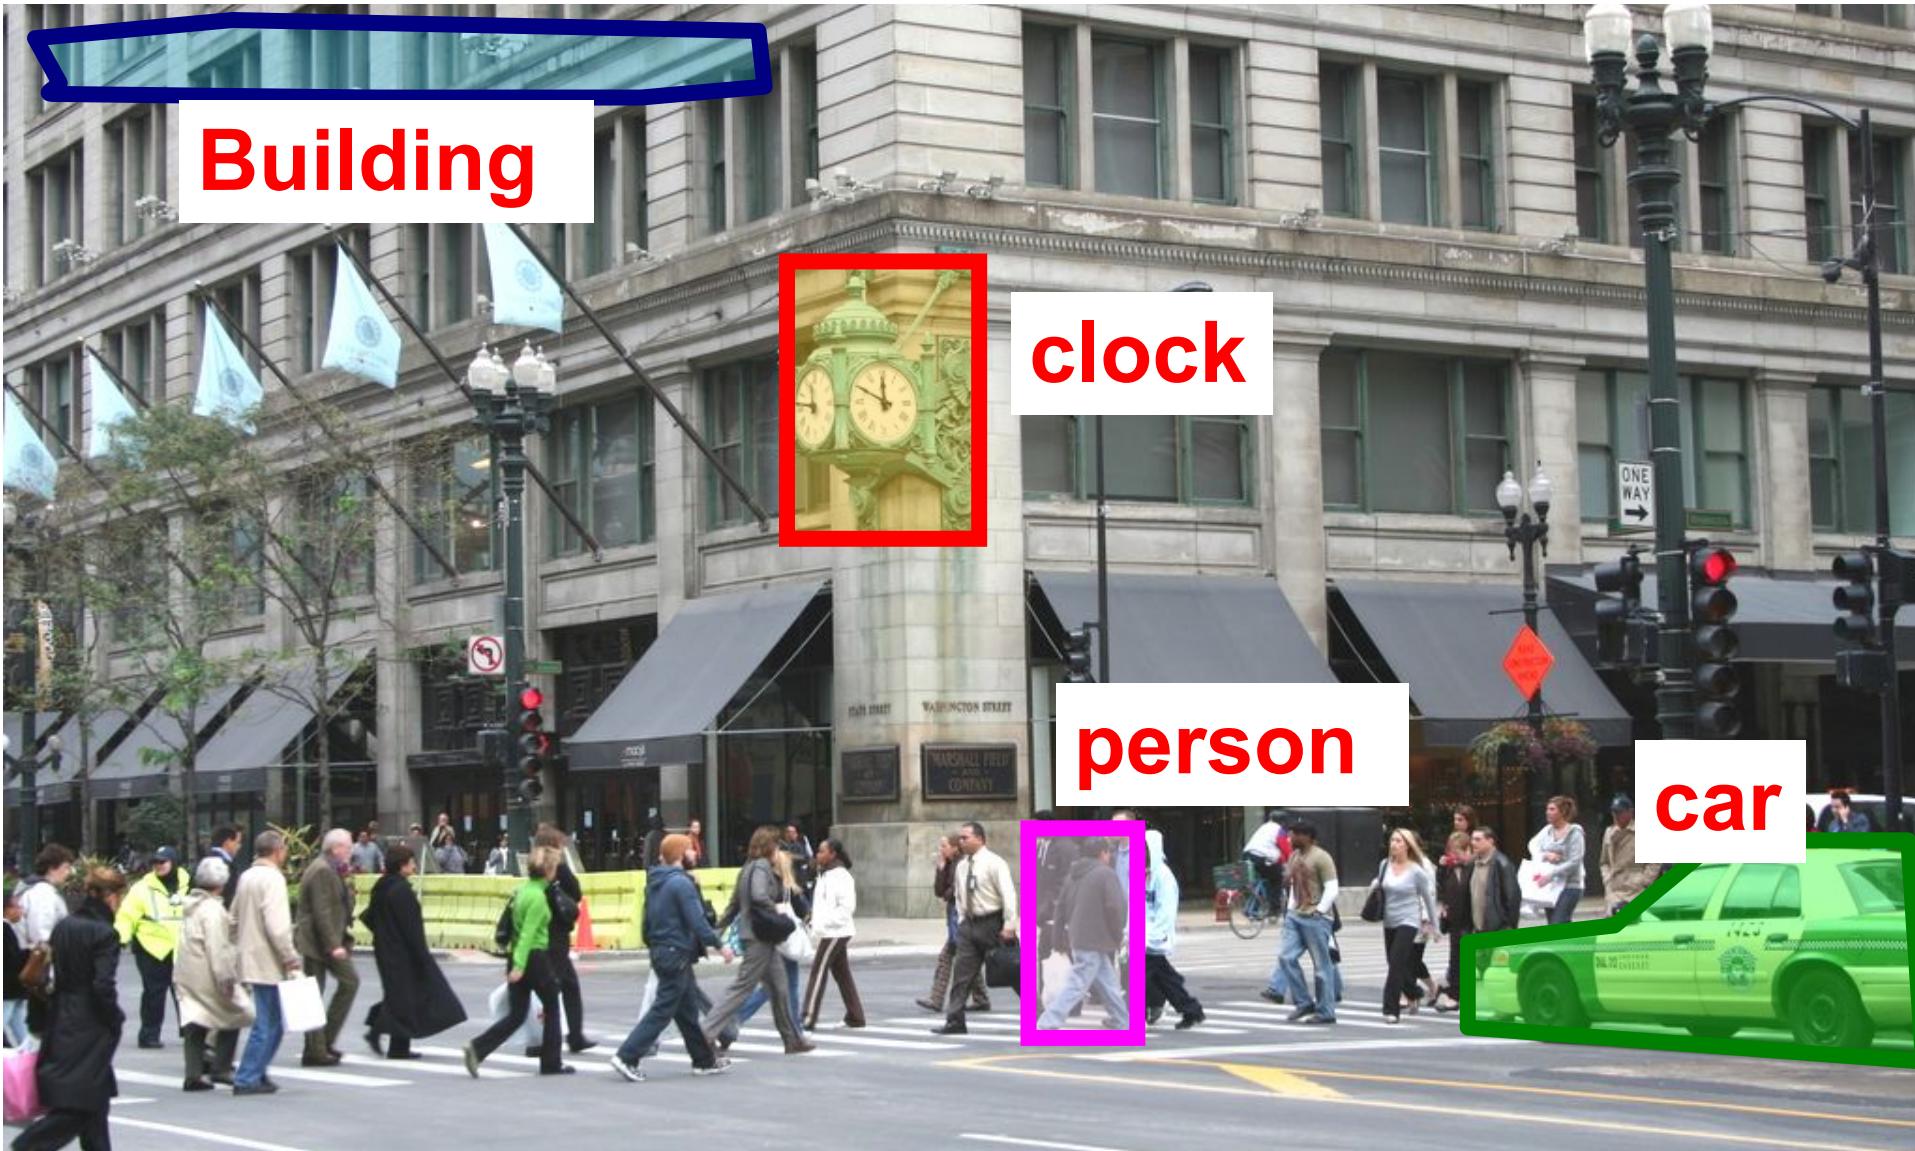
Applications: Image Search & Organizing photo collections



Detection: Does this image contain a car? [where?]



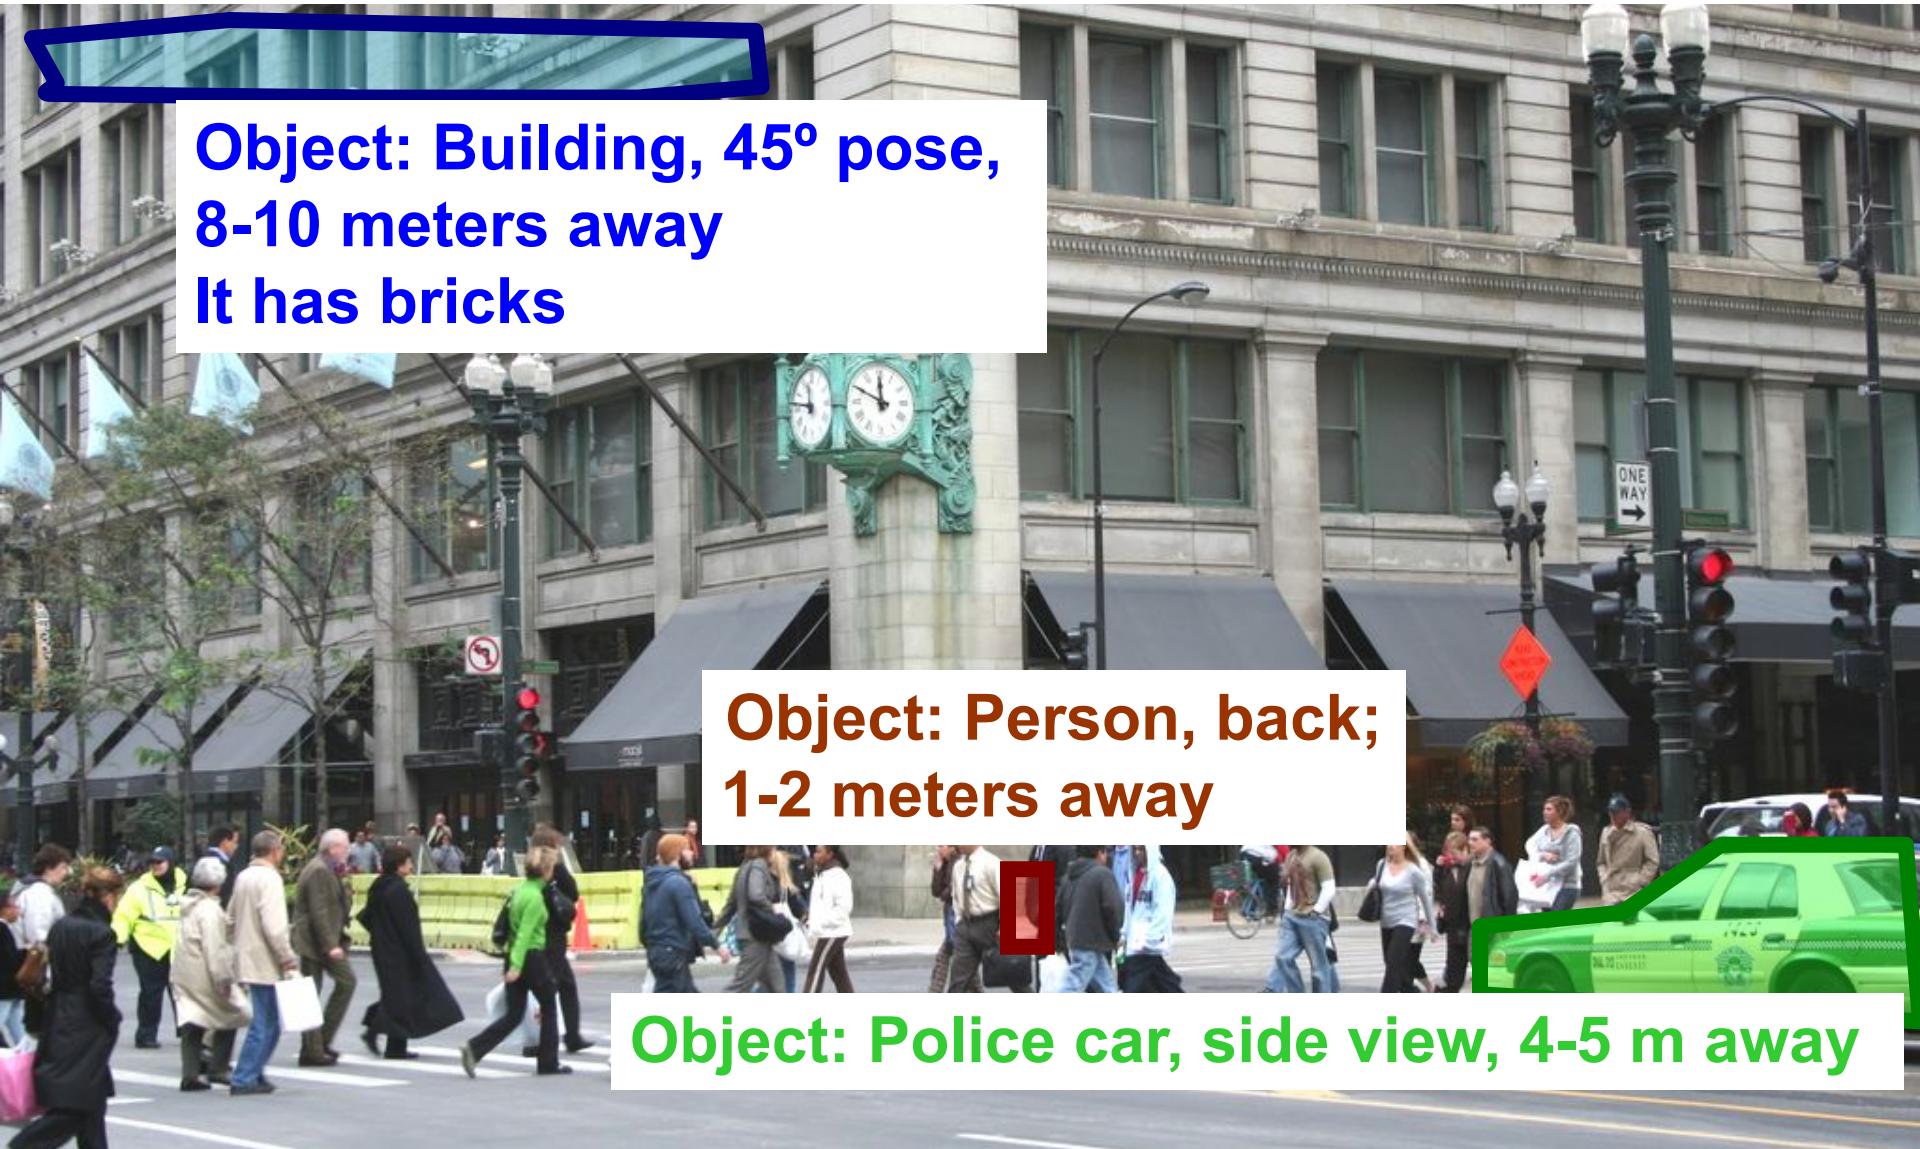
Detection: Which object does this image contain? [where?]



Detection: Accurate localization (segmentation)



Detection: Estimating object semantic & geometric attributes



Levels of recognition: Category-level vs instance-level

Does this image contain the Chicago Macy's building?



Categorization vs Single instance recognition

We have seen a form of single instance categorization already: **Where is the crunchy nut?**



Applications of computer vision



Recognizing landmarks
in mobile devices

Activity recognition: What are these people doing?



Visual Recognition

- Design algorithms that can:
 - Classify images or videos
 - Detect and localize objects
 - Estimate semantic and geometrical attributes
 - Classify human activities and events

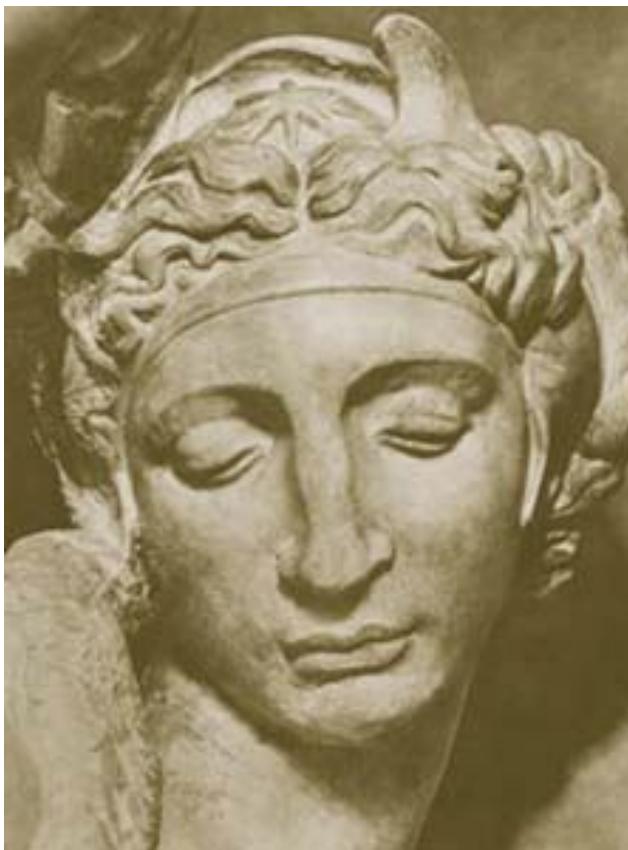
Why is this challenging?

How many
object
categories are
there?

~10,000 to 30,000



Challenges: viewpoint variation



Michelangelo 1475-1564

Challenges: illumination

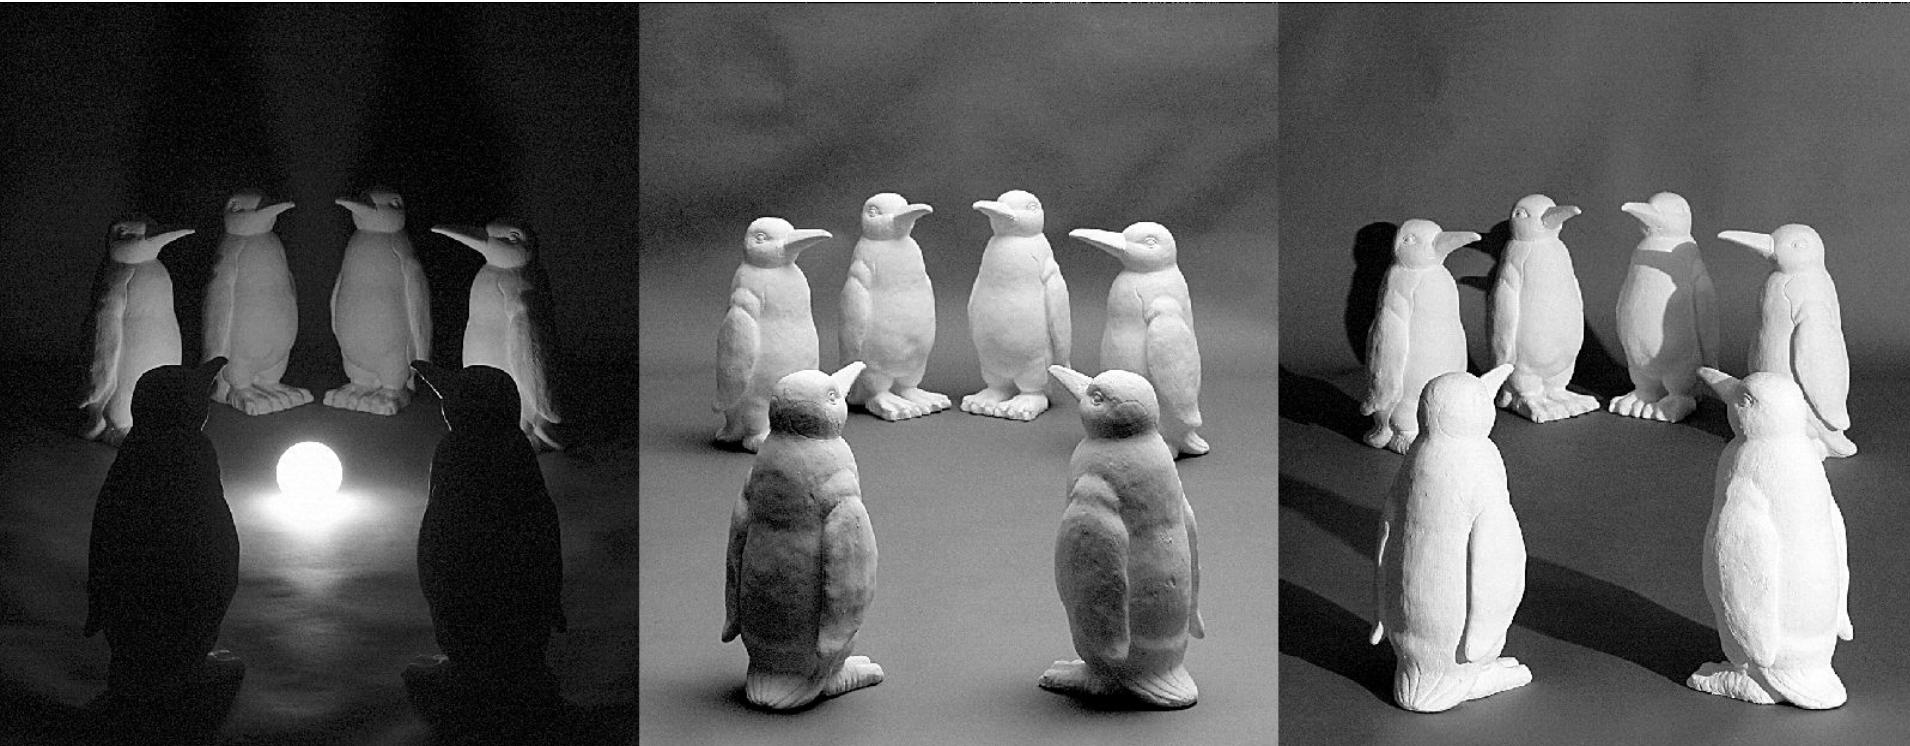


image credit: J. Koenderink

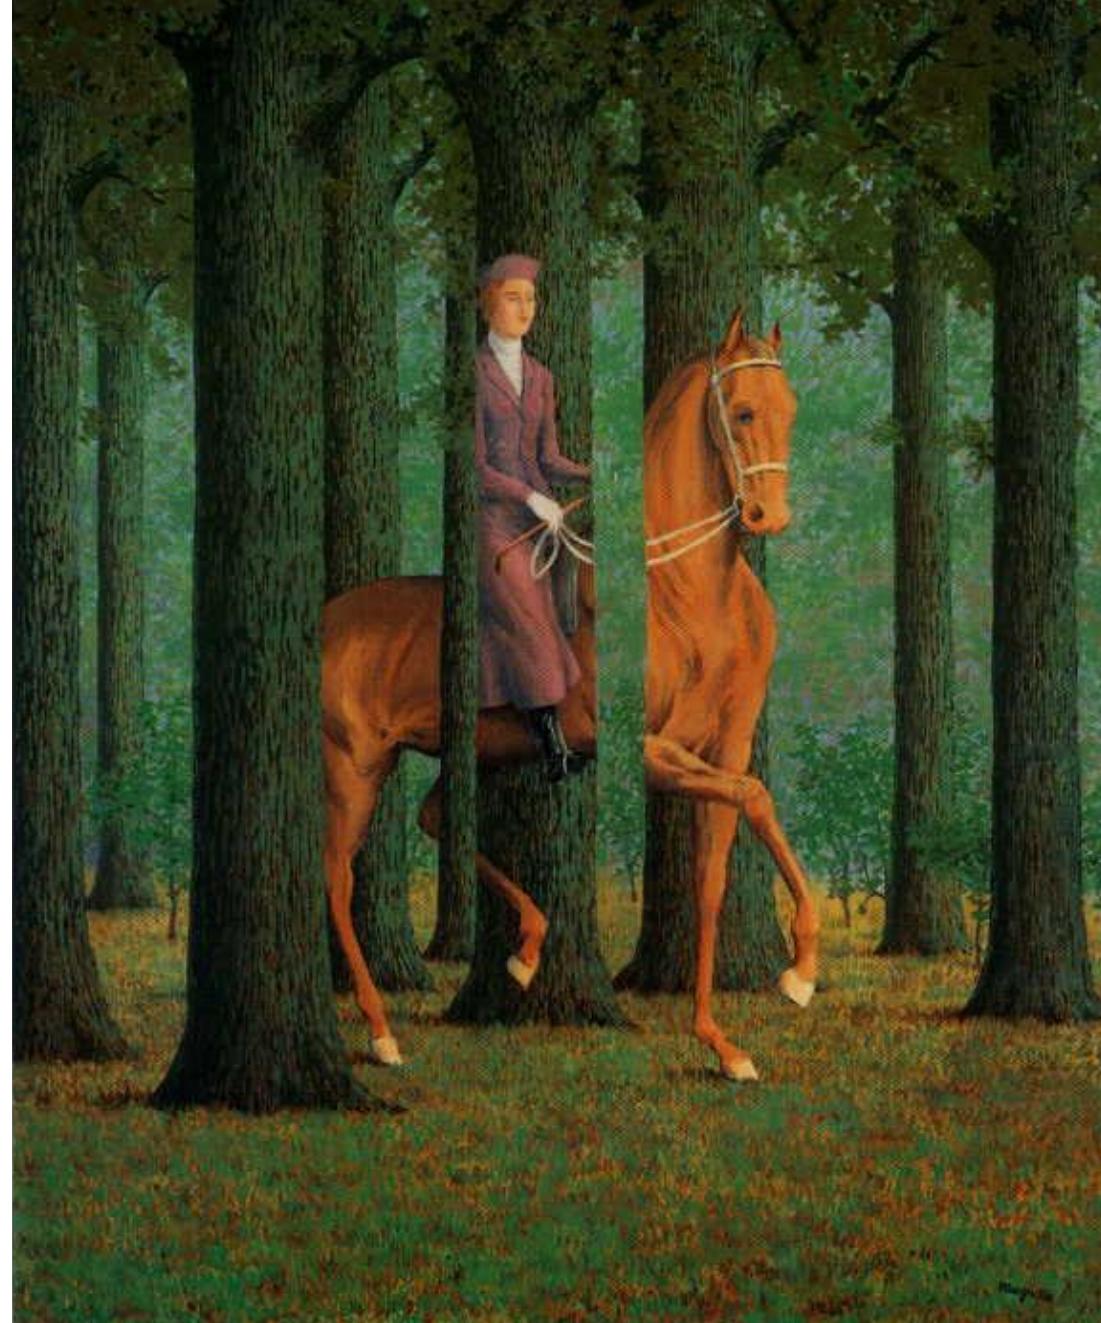
Challenges: scale



Challenges: deformation

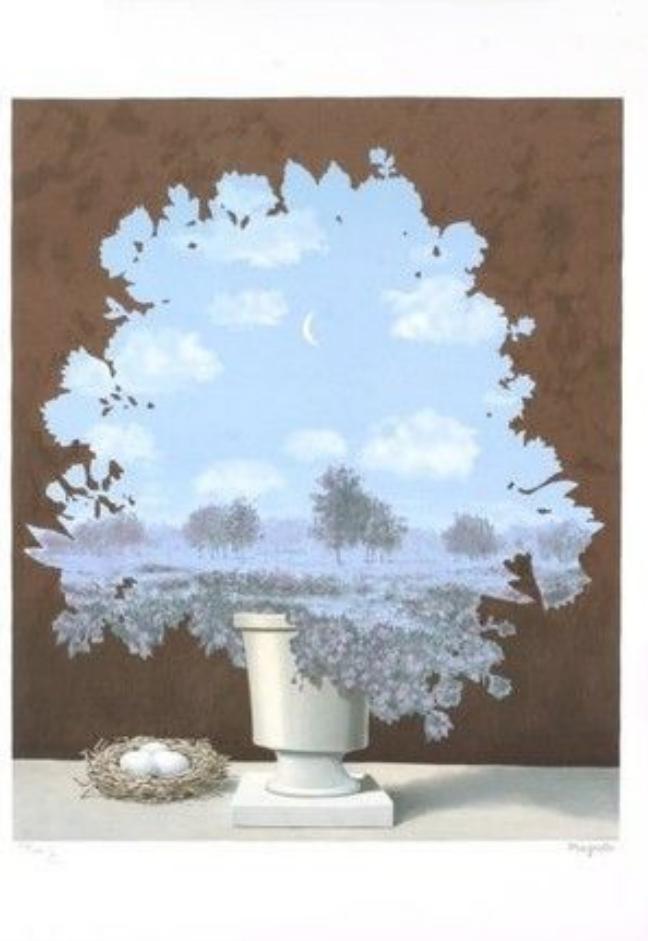


Challenges: occlusion



Magritte, 1957

Art Segway - Magritte

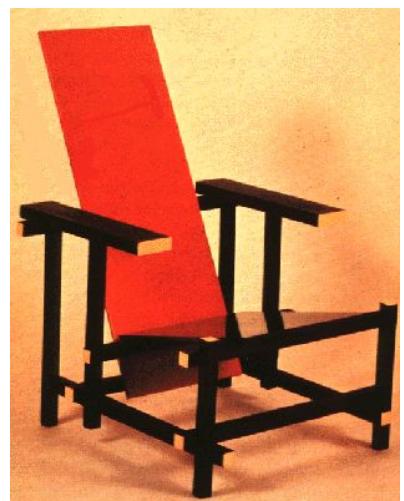


Challenges: background clutter

Kilmeny Niland. 1995



Challenges: intra-class variation



Today's agenda

- Introduction to recognition
- A object recognition pipeline
- Choosing the right features
- A training algorithm: KNN
- Testing an algorithm
- Challenges with kNN
- Dimensionality reduction

Object recognition: a classification framework

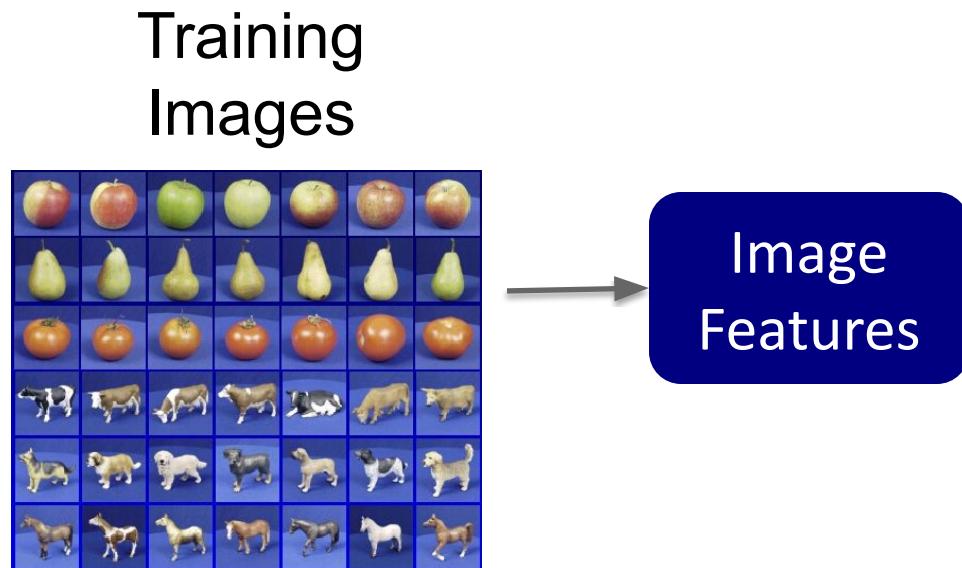
- Apply a prediction function to a feature representation of the image to get the desired output:

$f(\text{apple}) = \text{"apple"}$

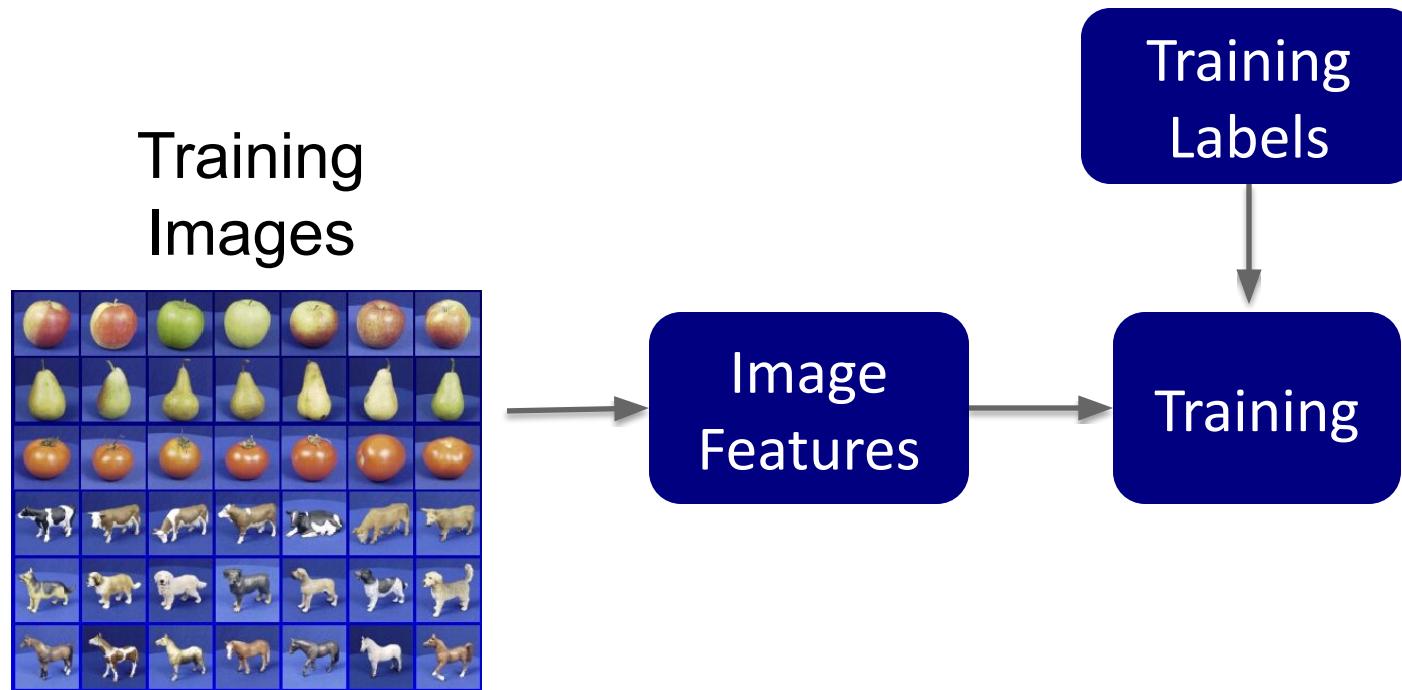
$f(\text{tomato}) = \text{"tomato"}$

$f(\text{cow}) = \text{"cow"}$

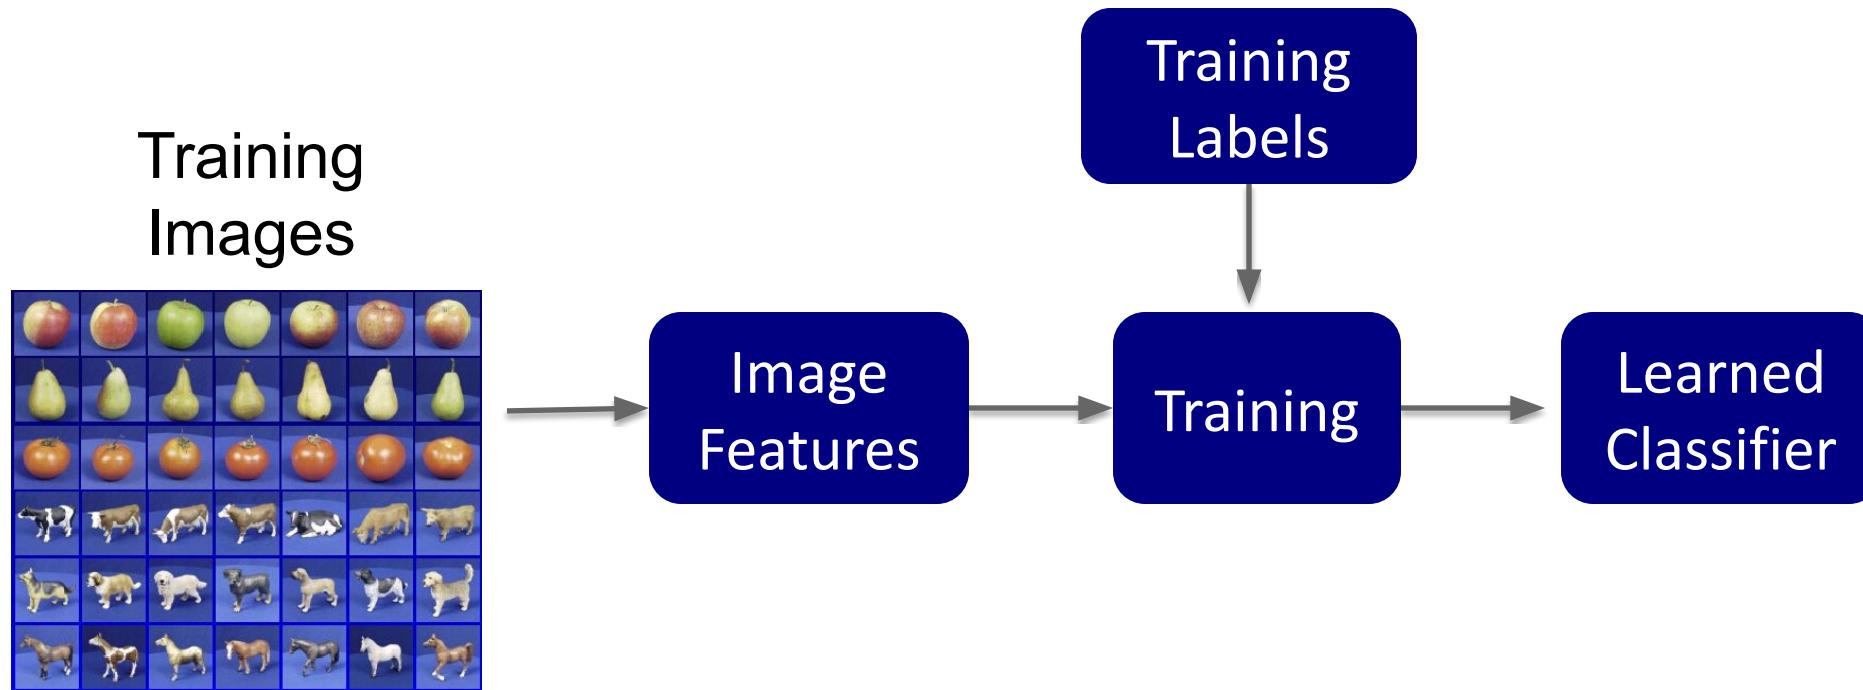
A simple pipeline - Training



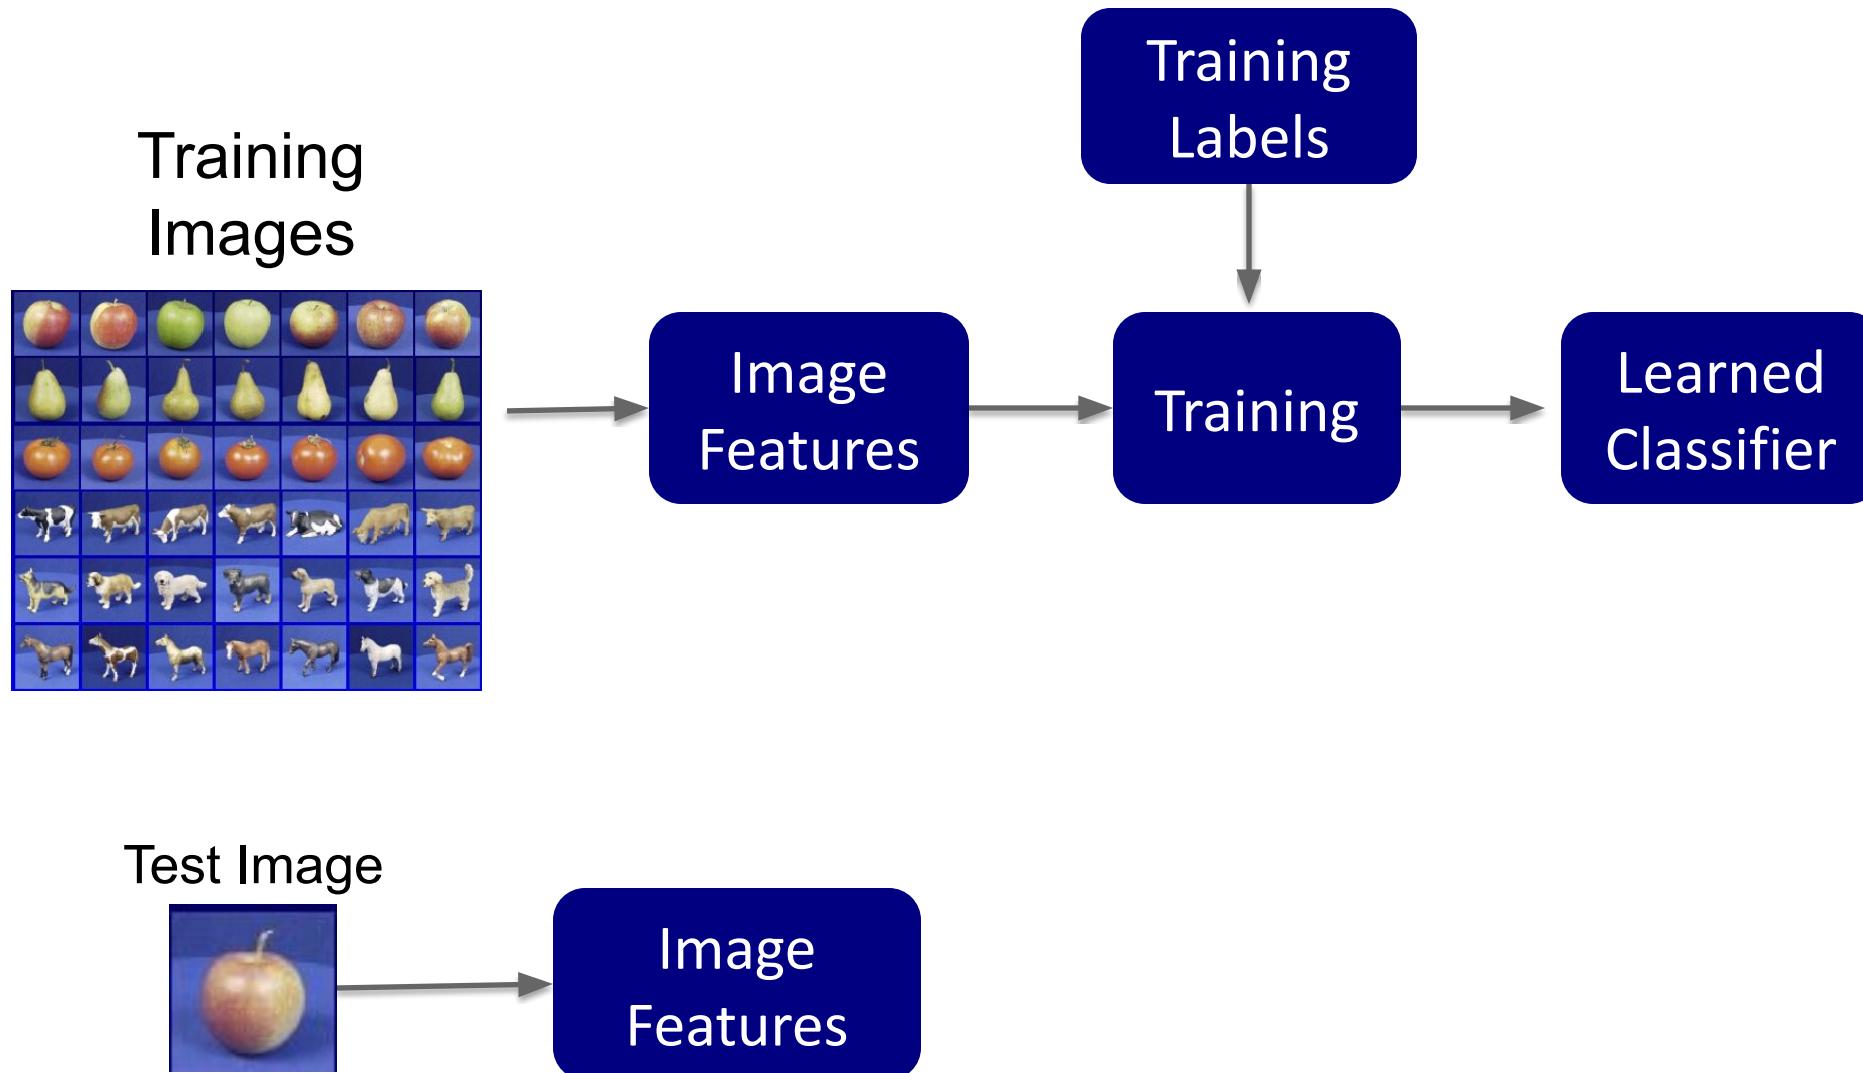
A simple pipeline - Training



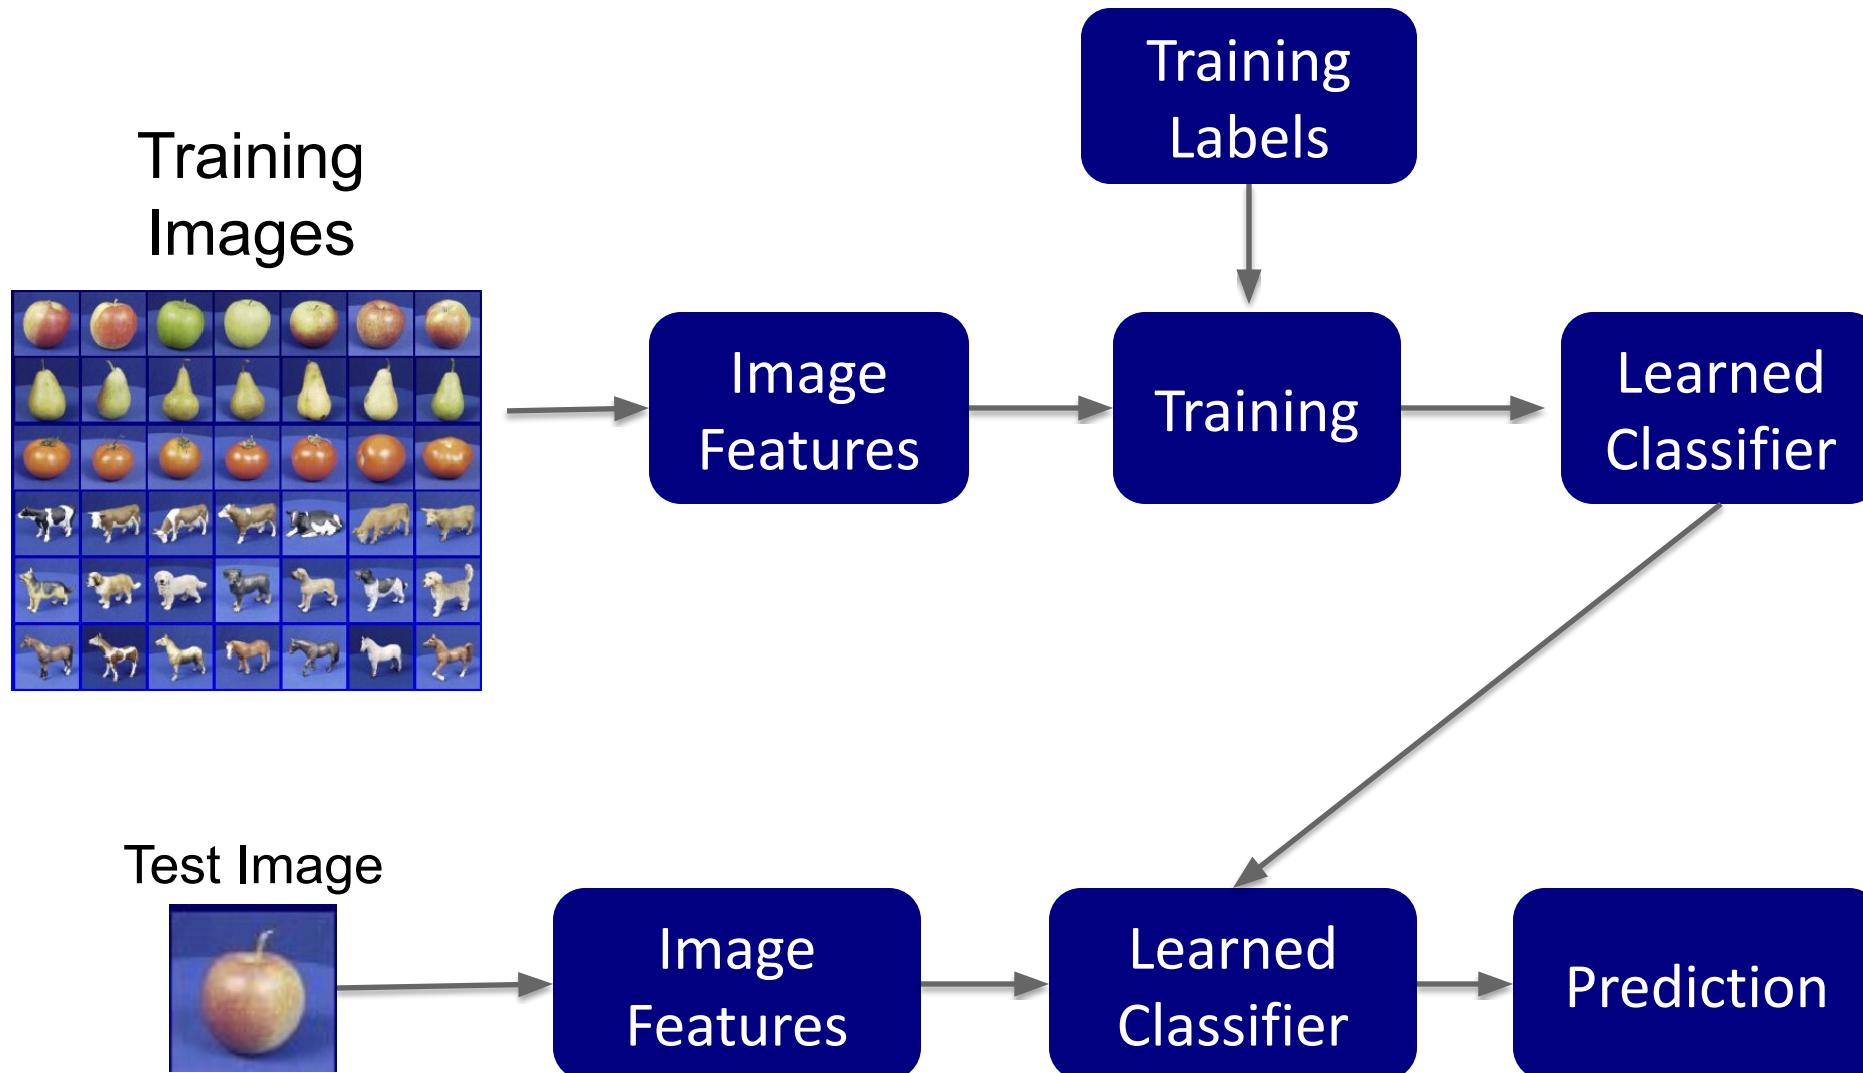
A simple pipeline - Training



A simple pipeline - Training



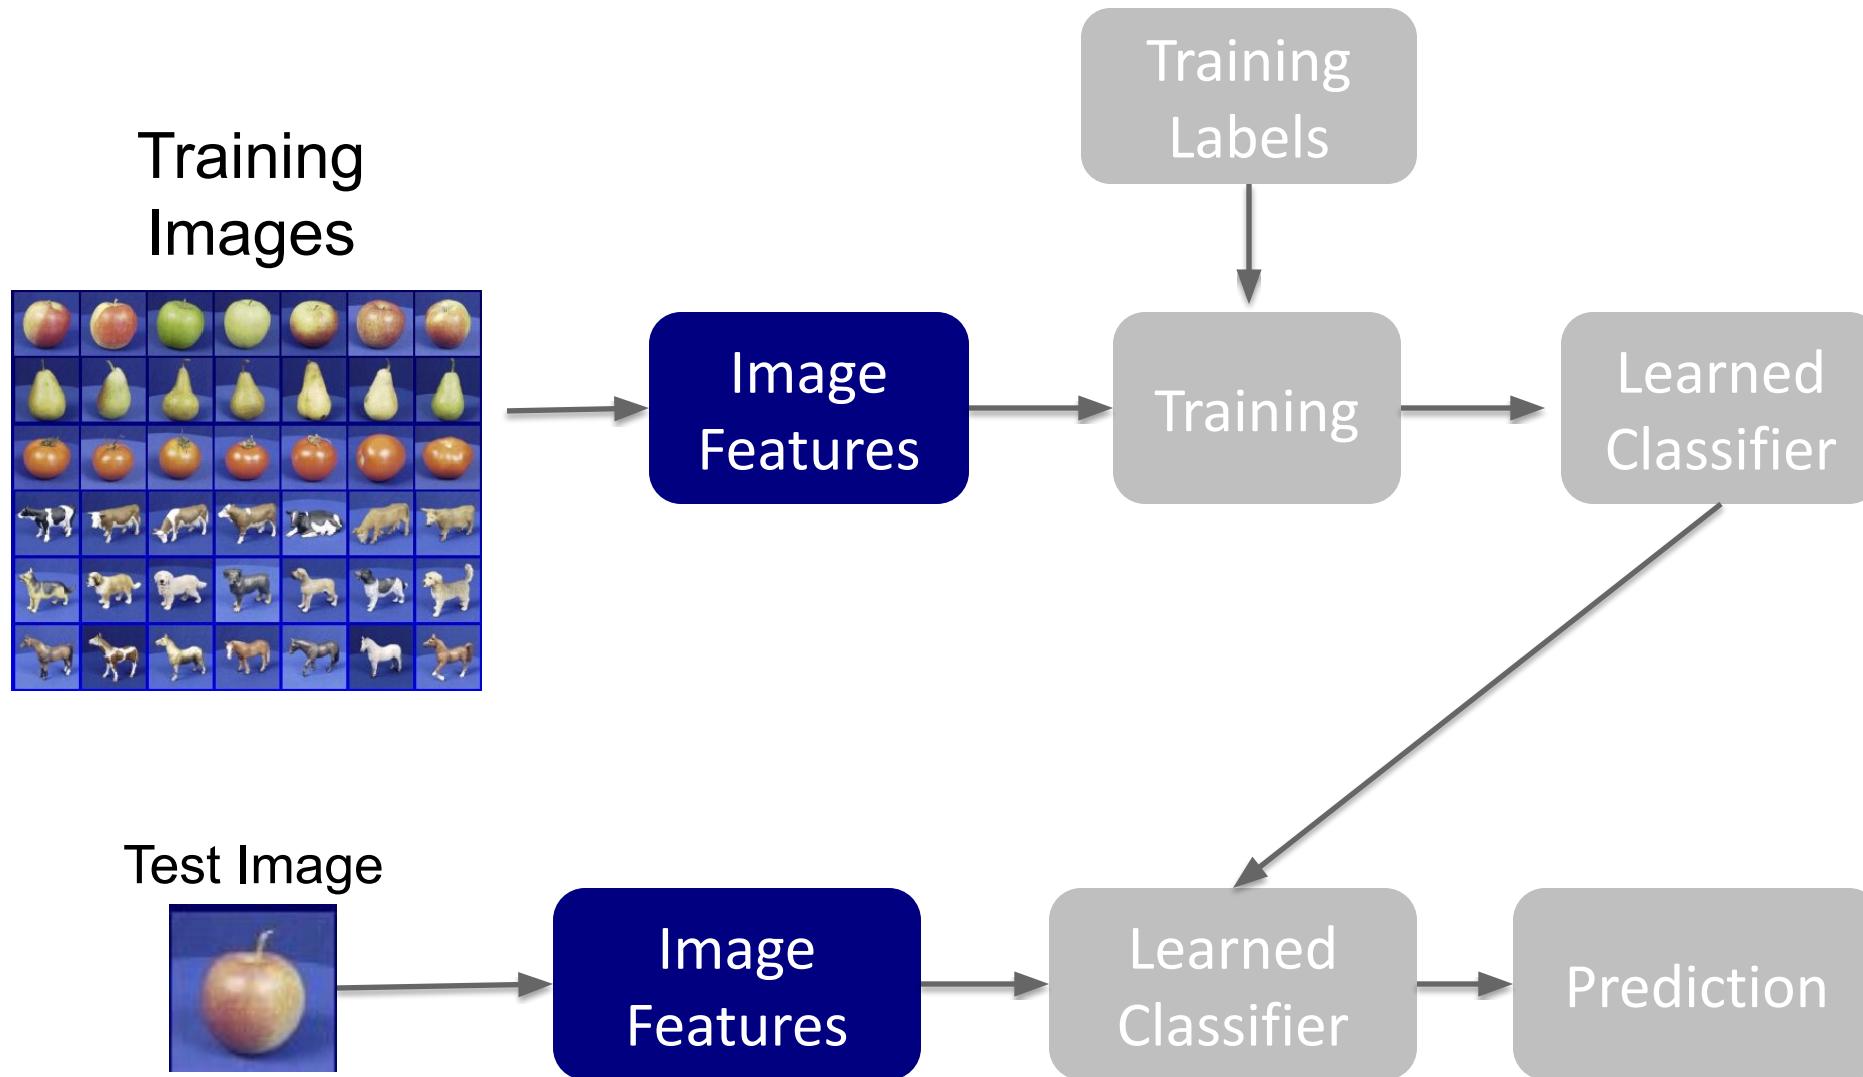
A simple pipeline - Training



What we will learn today?

- Introduction to recognition
- A object recognition pipeline
- **Choosing the right features**
- A training algorithm: KNN
- Testing an algorithm
- Challenges with kNN
- Dimensionality reduction

A simple pipeline - Training



Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ? | | | | | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | | | | | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |

✓ (global color counts don't change if the image shifts)

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ? | | | | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | | | | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |

✓ (if the *entire* image is uniformly scaled, the color distribution remains the same)

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ? | ? | | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |

- ✓ (rotating the entire image does not change overall color distribution)
- ✗ (appearance/colors can change if out-of-plane rotation reveals different surfaces)

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ? | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | | |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |

✗ (removing part of the image can significantly alter color histogram)

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ? | ? |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ |
| | | | | | | | |
| | | | | | | | |
| | | | | | | | |

✗ (shifts in illumination change color intensities/distribution)

✗ (noise directly alters pixel distribution)

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ |
| HoG | ? | ? | ? | ? | | | |
| | | | | | | | |
| | | | | | | | |

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ |
| HoG | ✗ | ✗ | ✗ | ✗ | | | |
| | | | | | | | |
| | | | | | | | |

✗ (local bins move)

✗ (needs re-computation at multiple scales)

✗ (oriented gradients are tied to an image grid)

✗ (same reason as the ^)

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ |
| HoG | ✗ | ✗ | ✗ | ✗ | ? | ? | ? |
| | | | | | | | |
| | | | | | | | |

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ |
| HoG | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ |
| | | | | | | | |
| | | | | | | | |

✗ (partial occlusion would result in no match)

✓ (gradients are more stable under monotonic intensity changes)

✗ (gradient orientations can be disrupted by significant noise)

Choices of features

| | Invariances | | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|--|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise | |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ | |
| HoG | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ | |
| SIFT | ? | ? | ? | ? | | | | |
| | | | | | | | | |

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ |
| HoG | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ |
| SIFT | ✓ | ✓ | ✓ | ✗ | | | |
| | | | | | | | |

- ✓ (keypoint-based, unaffected by shift)
- ✓ (built-in scale normalization)
- ✓ (SIFT normalizes orientation)
- ✗ (local keypoints might disappear if the object rotates)

Choices of features

| | Invariances | | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|--|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise | |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ | |
| HoG | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ | |
| SIFT | ✓ | ✓ | ✓ | ✗ | ? | ? | ? | |
| | | | | | | | | |

Choices of features

| | Invariances | | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|--|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise | |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ | |
| HoG | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ | |
| SIFT | ✓ | ✓ | ✓ | ✗ | ✓ | ✓ | ✓ | |
| Deep learning | | | | | | | | |

- ✓ (local keypoints can still match if some are visible)
- ✓ (gradient-based + normalization)
- ✓ (SIFT is relatively robust to moderate noise)

Choices of features

| | Invariances | | | | | | | |
|---------------|-------------|-------|---|-----------------------------|----------------------|--------------|-------------------|--|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise | |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ | |
| HoG | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ | |
| SIFT | ✓ | ✓ | ✓ | ✗ | ✓ | ✓ | ✓ | |
| Deep learning | ? | ? | ? | ? | | | | |

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|---------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ |
| HoG | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ |
| SIFT | ✓ | ✓ | ✓ | ✗ | ✓ | ✓ | ✓ |
| Deep learning | usually | usually | usually | ✗ | | | |

~ Deep learning features are **usually** invariant to translation, scale, and planar rotation if the training data has these translations. It is not invariant to other rotations.

Aside: ImageNet has objects centered in the middle of images. So models trained on ImageNet are not translation or scale invariant.

Choices of features

| | Invariances | | | | | | | |
|---------------|-------------|---------|---|-----------------------------|----------------------|--------------|-------------------|--|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise | |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ | |
| HoG | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ | |
| SIFT | ✓ | ✓ | ✓ | ✗ | ✓ | ✓ | ✓ | |
| Deep learning | usually | usually | usually | ✗ | ? | ? | ? | |

Choices of features

| | Invariances | | | | | | |
|---------------|-------------|---------|---|-----------------------------|----------------------|--------------|-------------------|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ |
| HoG | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ |
| SIFT | ✓ | ✓ | ✓ | ✗ | ✓ | ✓ | ✓ |
| Deep learning | usually | usually | usually | ✗ | ✗ | ✗ | ✓ |

- ✗ (standard CNNs are not strictly occlusion-invariant; partial robustness depends on training)
✓ (can learn robustness if trained on varied lighting)
✓ (CNNs can learn to be noise-robust with proper training)

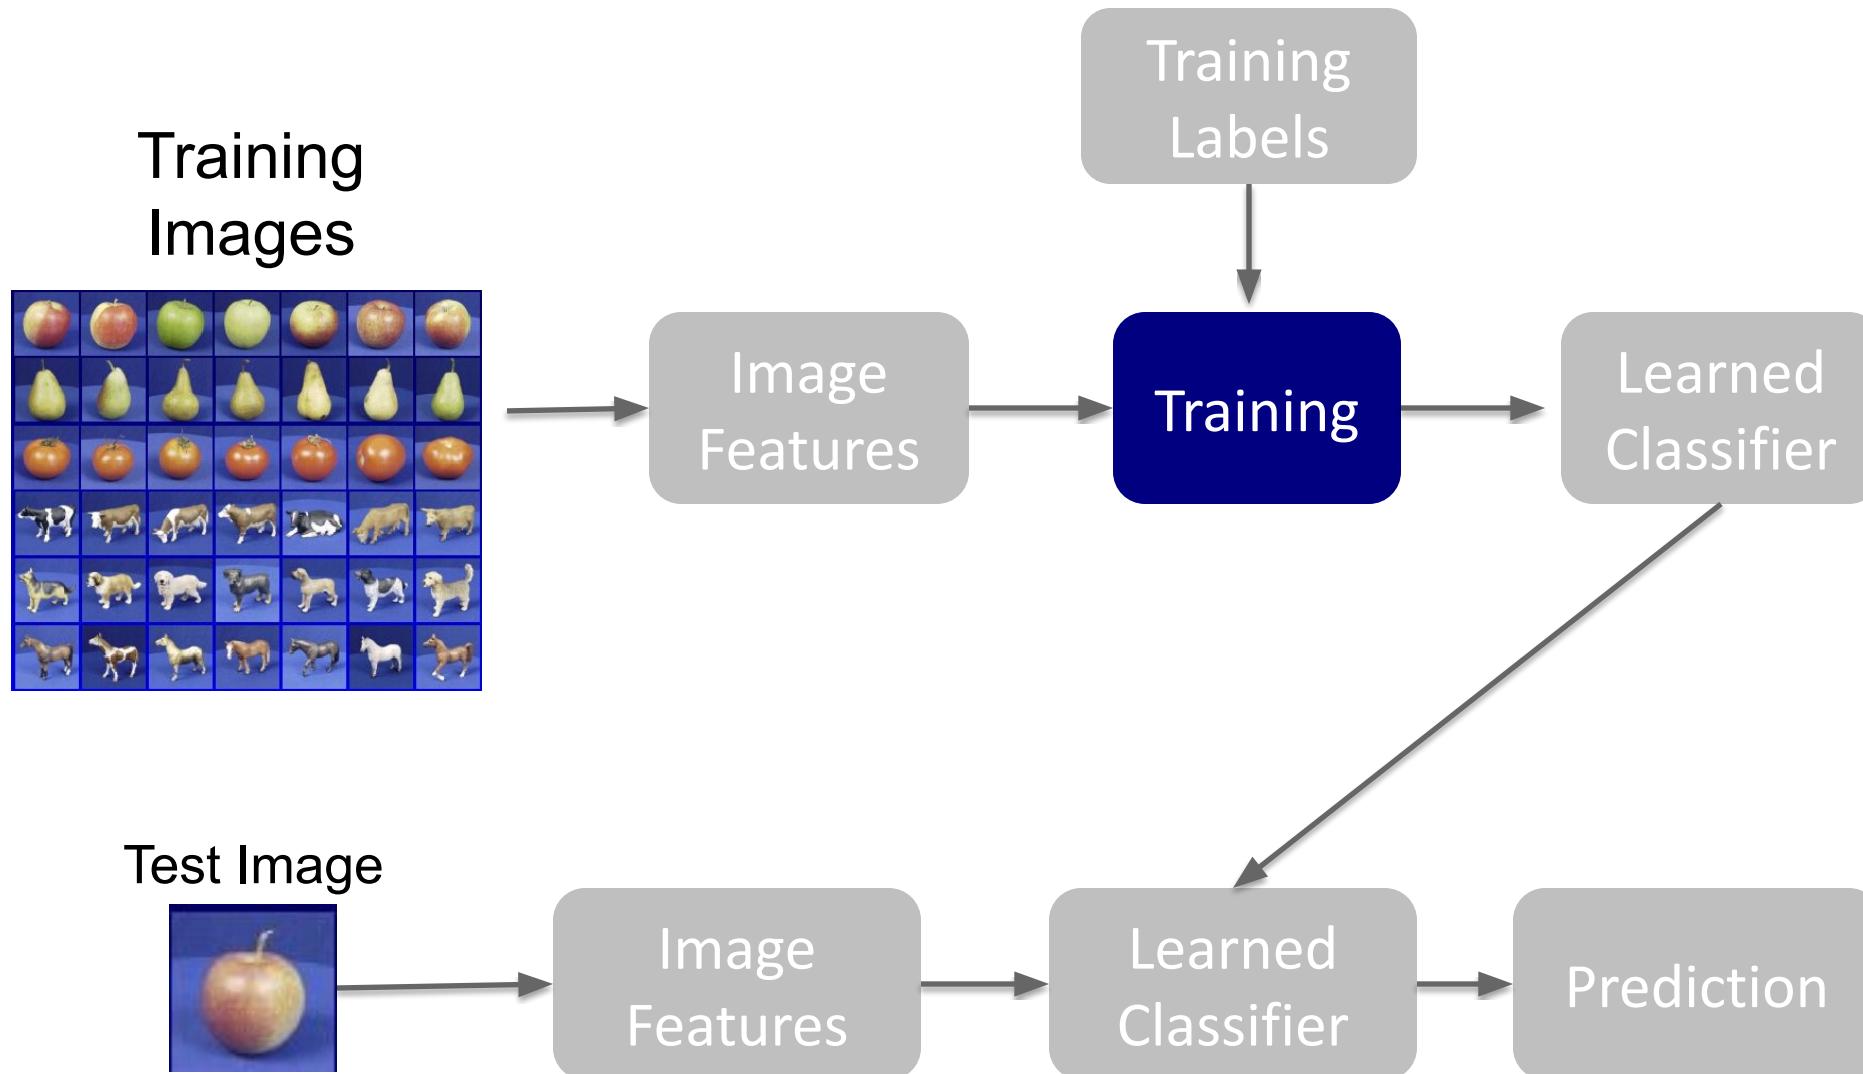
So, which features should we choose?

| | Invariances | | | | | | | |
|---------------|-------------|---------|---|-----------------------------|----------------------|--------------|-------------------|--|
| | Translation | Scale | Rotation (relative to camera plane) | Rotation (unconstrained) | Partial Occlusion | Illumination | Gaussian Noise | |
| RGB-histogram | ✓ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ | |
| HoG | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ | |
| SIFT | ✓ | ✓ | ✓ | ✗ | ✓ | ✓ | ✓ | |
| Deep learning | usually | usually | usually | ✗ | ✗ | ✓ | ✓ | |

What we will learn today?

- Introduction to recognition
- A simple Object Recognition pipeline
- Choosing the right features
- **A training algorithm: KNN**
- Testing an algorithm
- Challenges with kNN
- Dimensionality reduction

A simple pipeline - Training



Many classifiers to choose from

- **K-nearest neighbor**
- SVM
- Neural networks
- Naïve Bayes
- Bayesian network
- Logistic regression
- Randomized Forests
- Boosted Decision Trees
- RBMs
- Etc.

Which is the best one?

Learning a classifier to map inputs to outputs

$$y = f(\mathbf{x})$$

The diagram illustrates the mapping process. At the top is the equation $y = f(\mathbf{x})$. Below it, three blue arrows point upwards from labels to their corresponding components: a vertical arrow labeled "output" points to y ; another vertical arrow labeled "prediction function" points to $f(\mathbf{x})$; and a diagonal arrow labeled "Image feature" points to \mathbf{x} .

- **Training:** given a *training set* of labeled examples $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$, estimate the prediction function f by minimizing the prediction error on the training set
- **Testing:** apply f to a never before seen *test example* \mathbf{x} and output the predicted value $y = f(\mathbf{x})$

An example training dataset



Training set (labels known)

Apples

Pear

Tomatos

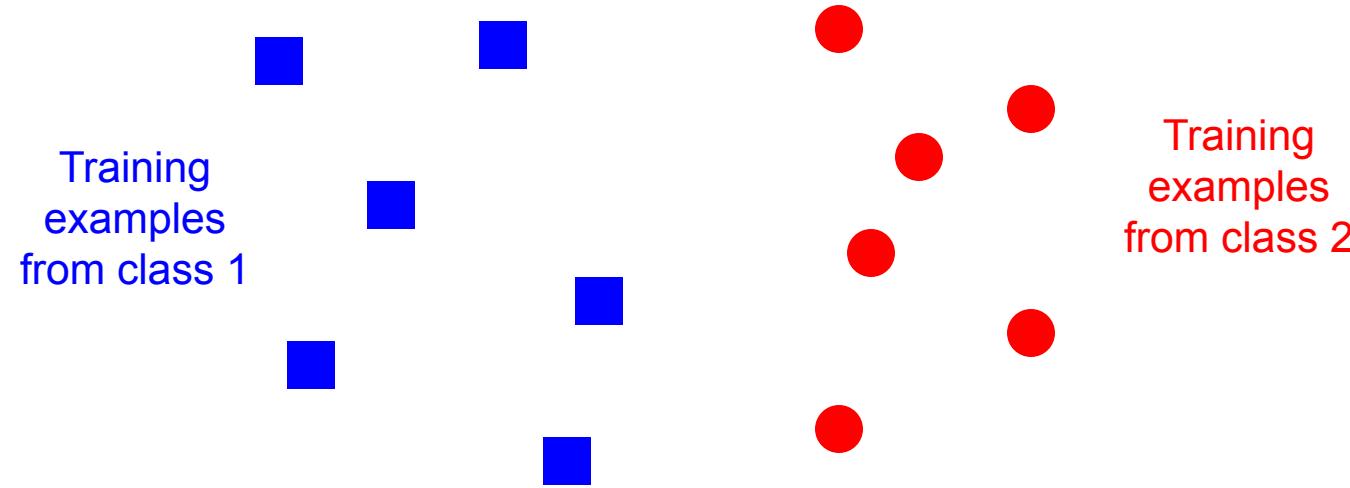
Cow

Dog

Horse

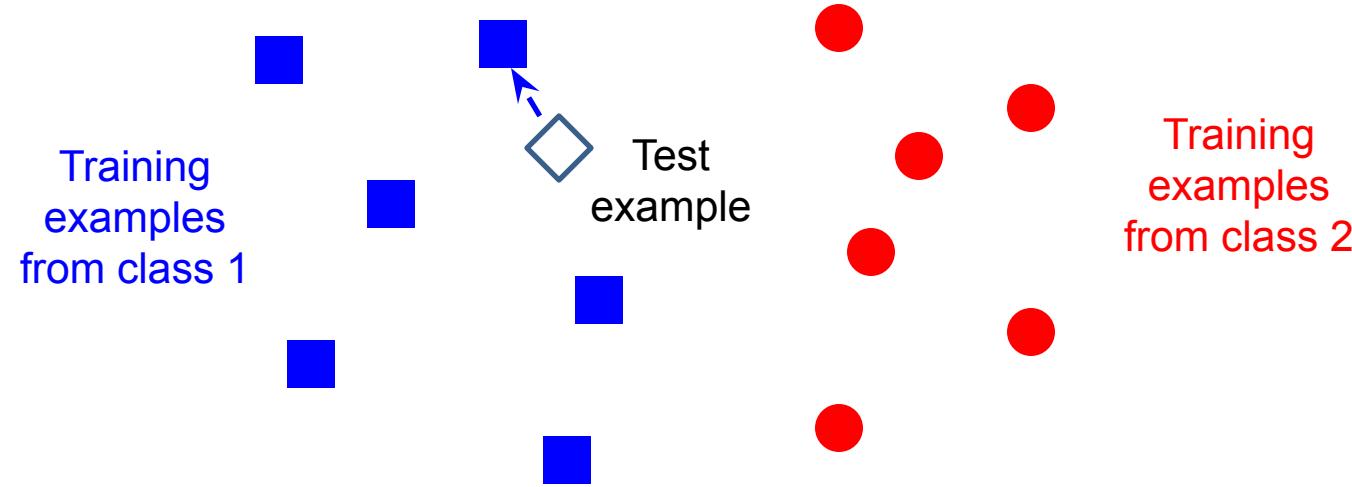
For kNN classifier,
training simply
means to store all
training data.

A stored training set



Slide credit: L. Lazebnik

During testing, we assign the label of the nearest neighbor in feature space

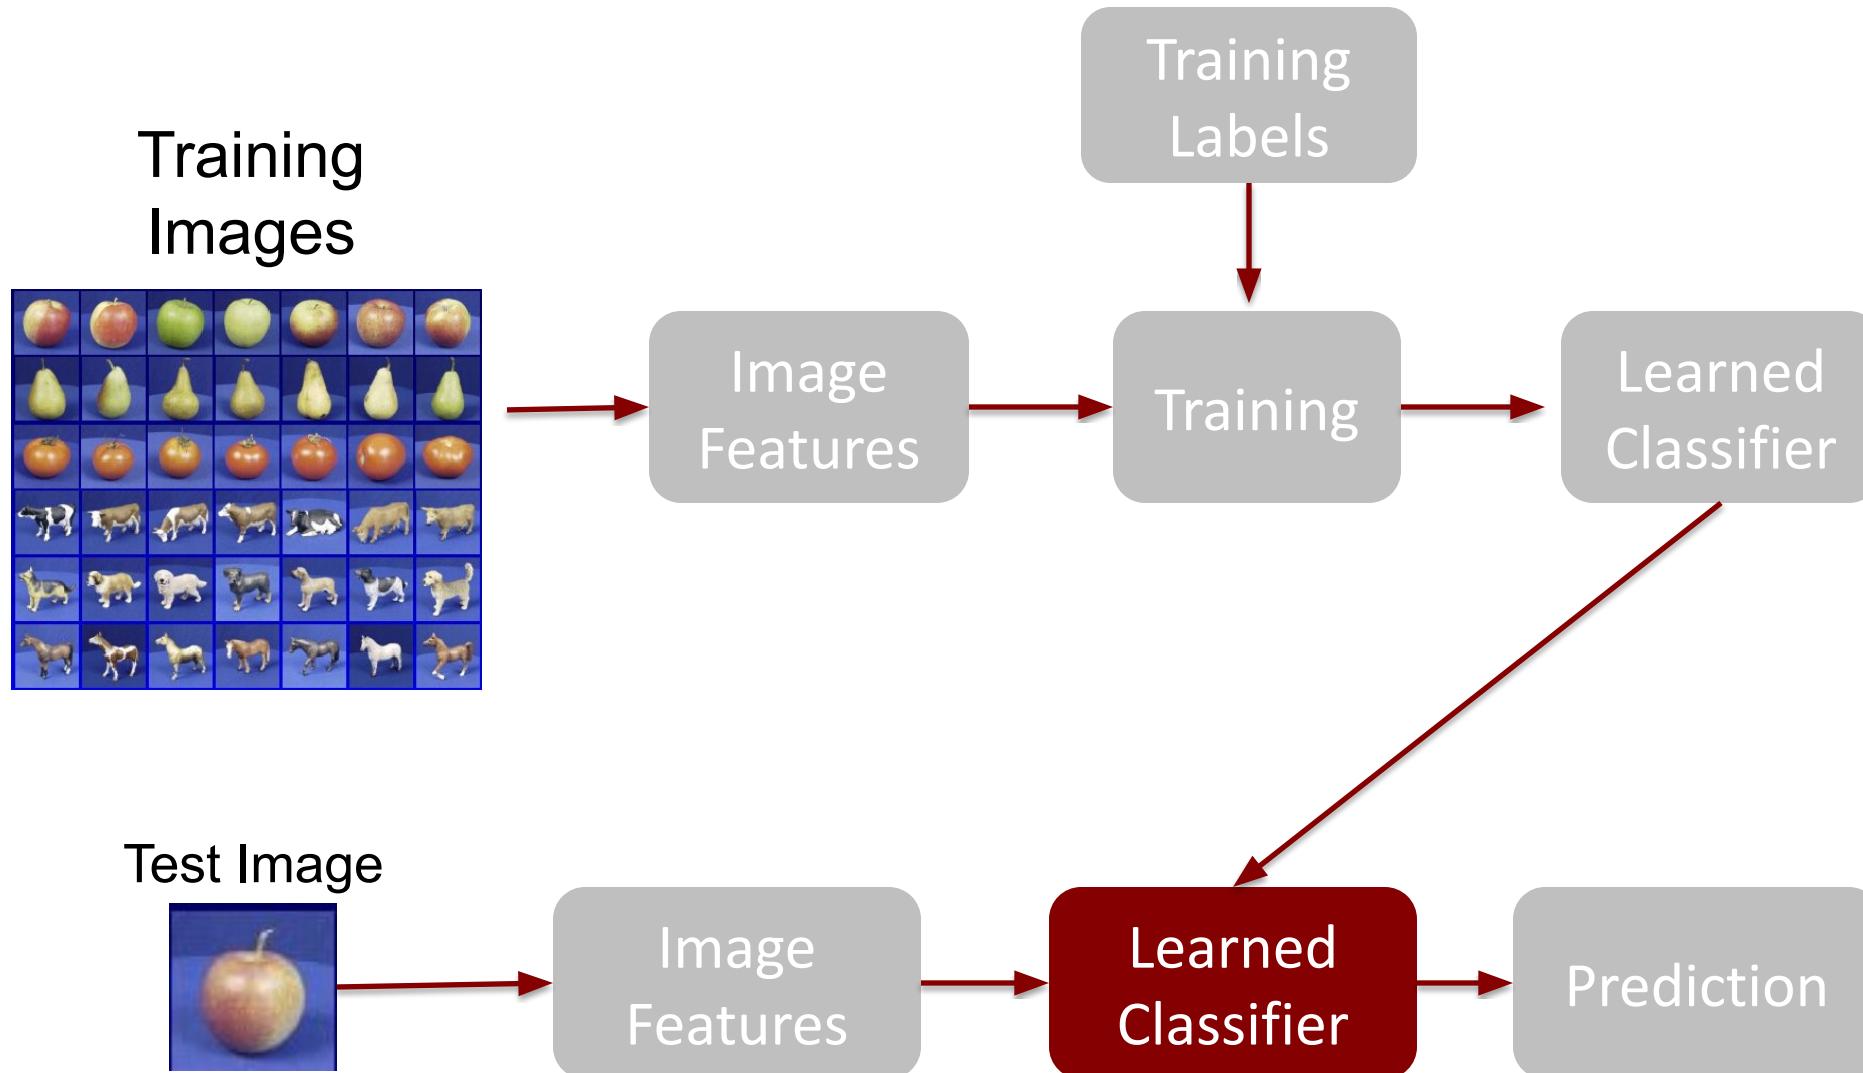


Slide credit: L. Lazebnik

What we will learn today?

- Introduction to recognition
- A simple Object Recognition pipeline
- Choosing the right features
- A training algorithm: kNN
- **Testing an algorithm**
- Challenges with kNN
- Dimensionality reduction

A simple pipeline - Training



Generalization



Training set (labels known)

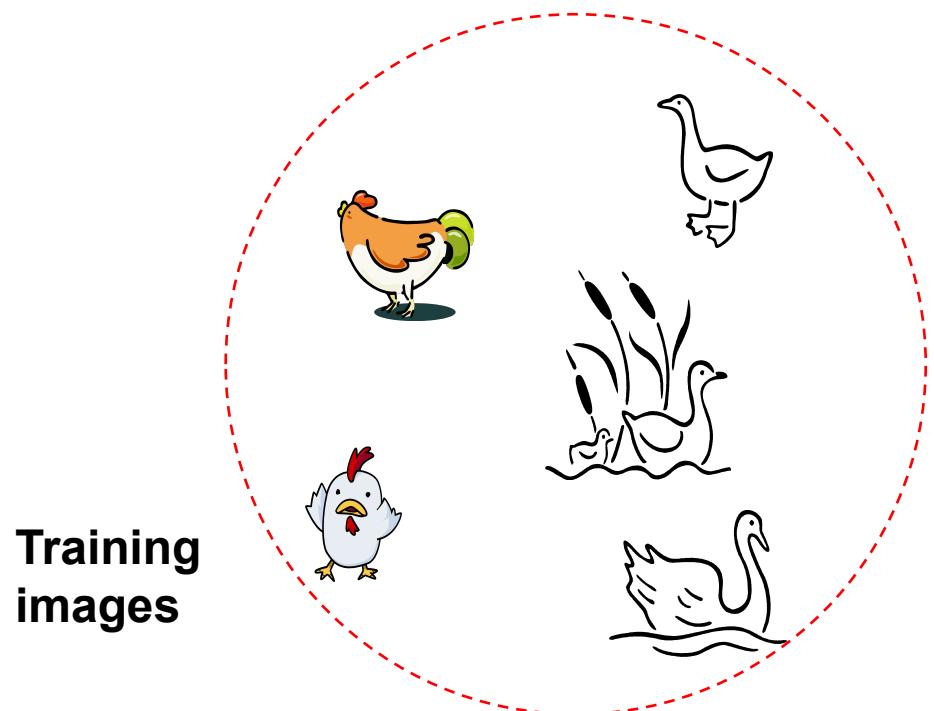


Test set (labels unknown)

- How well does a learned model generalize from the data it was trained on to a new test set?

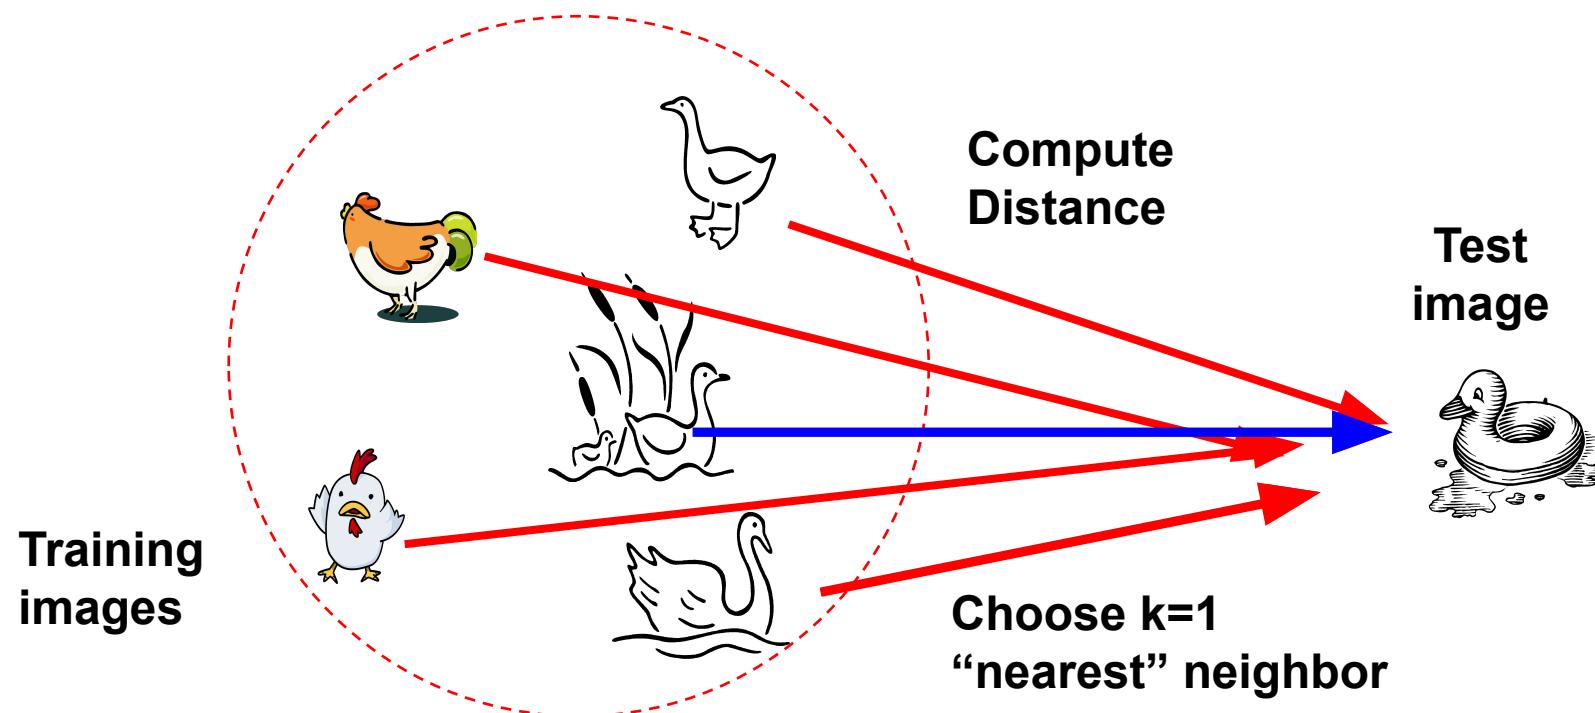
Intuition for Nearest Neighbor Classifier

Given a training dataset, simply store each image's features and their corresponding label.



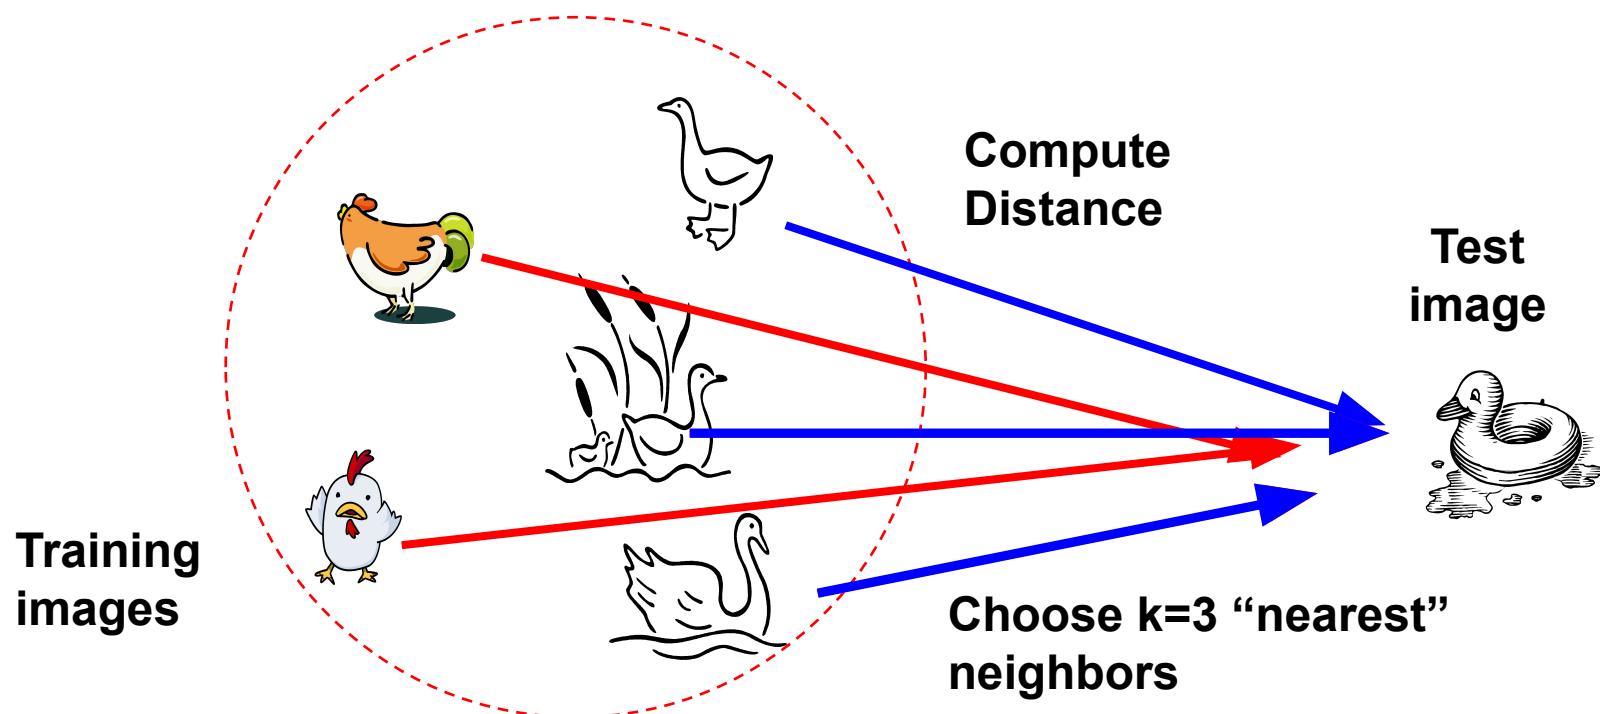
Intuition for Nearest Neighbor Classifier

Given a training dataset, simply store each image's features and their corresponding label.



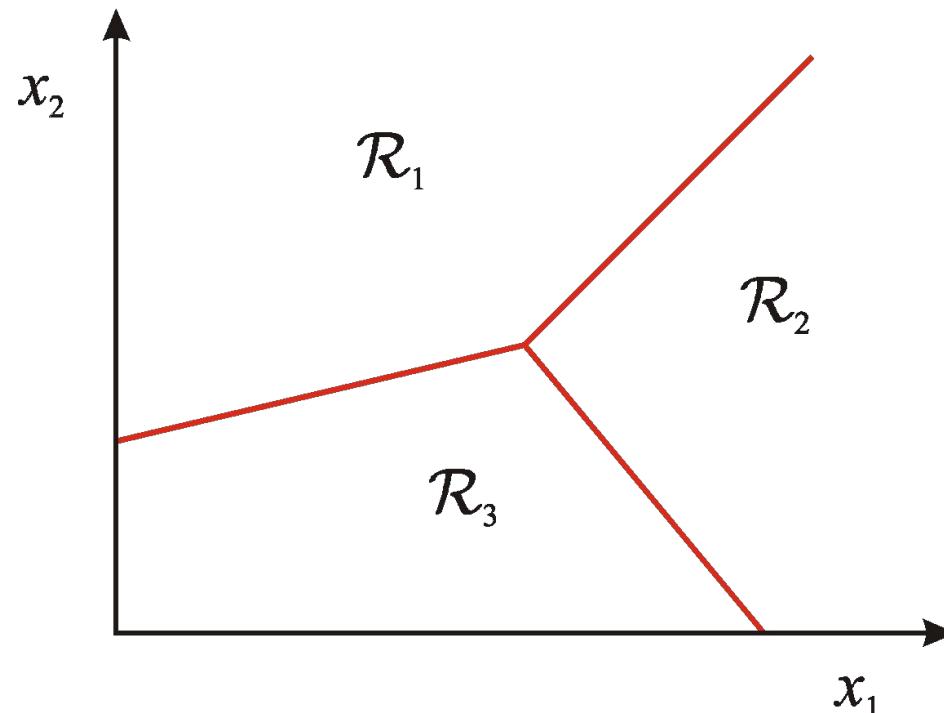
Nearest Neighbor Classifier

- Assign label of majority of K=3 nearest neighbors



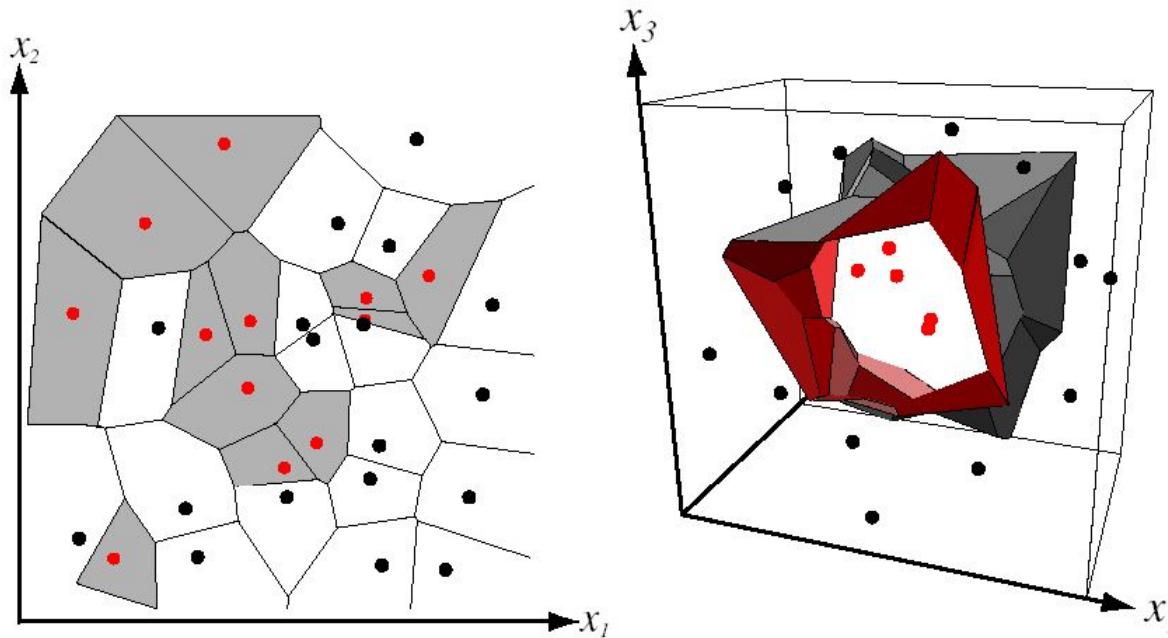
Classification

- Assign input vector to one of many classes (categories)
- **Geometric interpretation** of classifiers: A classifier divides input space into *decision regions* separated by *decision boundaries*



Nearest Neighbor Classifier

- Assign label of nearest training data point to each test data point



from Duda et al.

Partitioning of feature space
for two-category 2D and 3D data

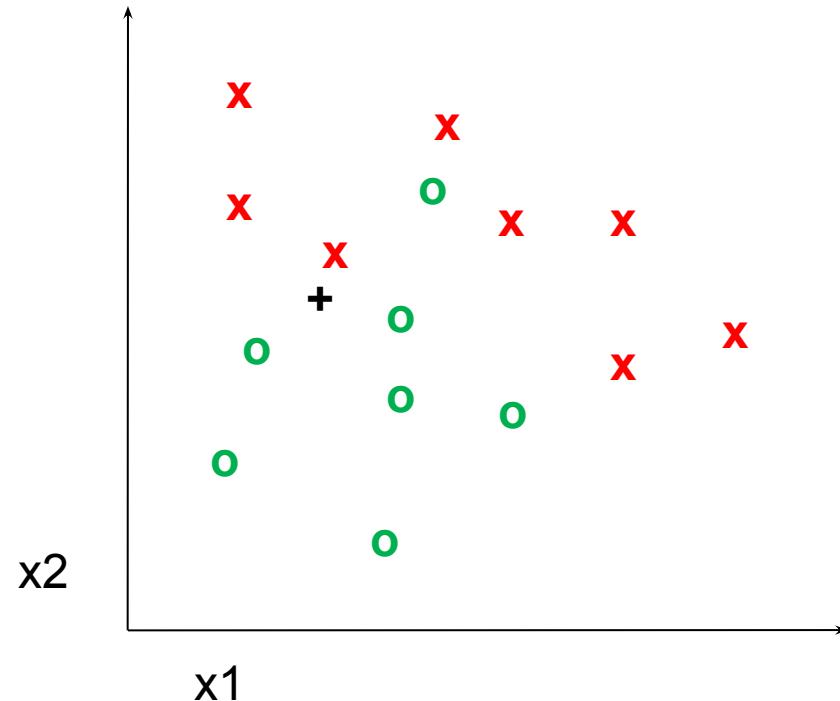
How do we find the nearest neighbors in feature space?

Distance measure (same as the ones from segmentation)

Euclidean:

$$Dist(X^n, X^m) = \sqrt{\sum_{i=1}^D (X_i^n - X_i^m)^2}$$

Where X^n and X^m are the n-th and m-th data points



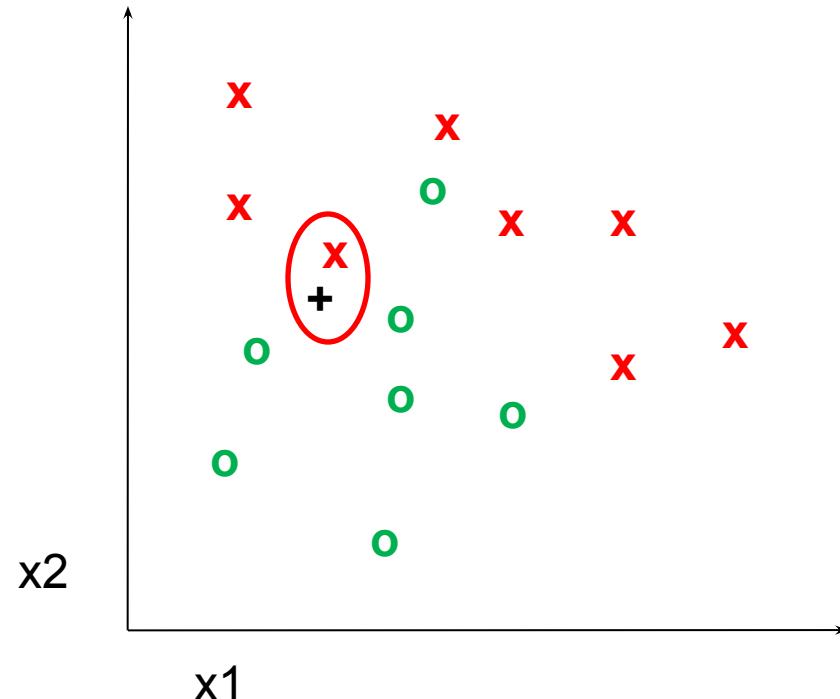
1-nearest neighbor

Distance measure (same as the ones from segmentation)

Euclidean:

$$Dist(X^n, X^m) = \sqrt{\sum_{i=1}^D (X_i^n - X_i^m)^2}$$

Where X^n and X^m are the n-th and m-th data points



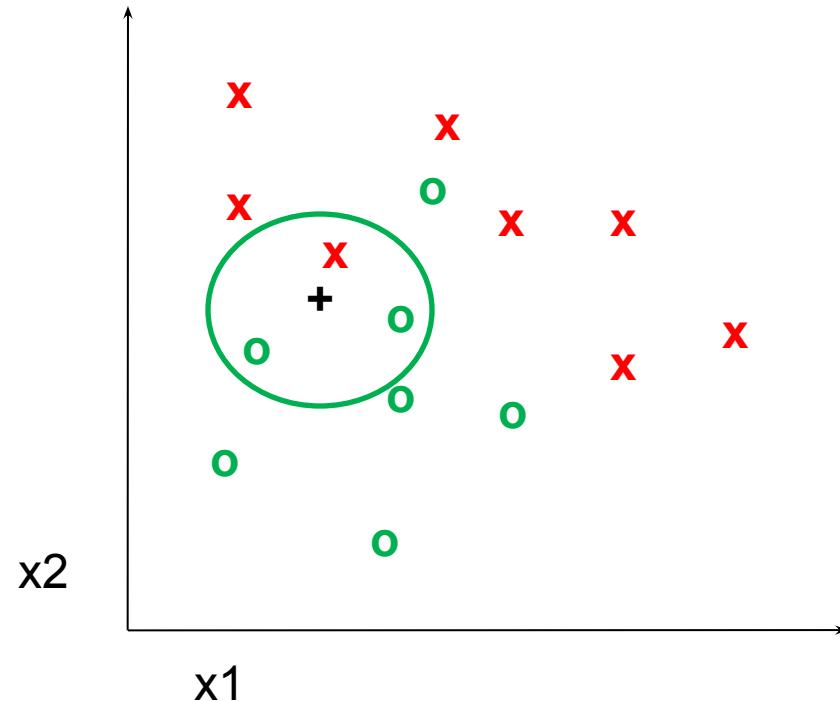
3-nearest neighbor

Distance measure (same as the ones from segmentation)

Euclidean:

$$Dist(X^n, X^m) = \sqrt{\sum_{i=1}^D (X_i^n - X_i^m)^2}$$

Where X^n and X^m are the n-th and m-th data points



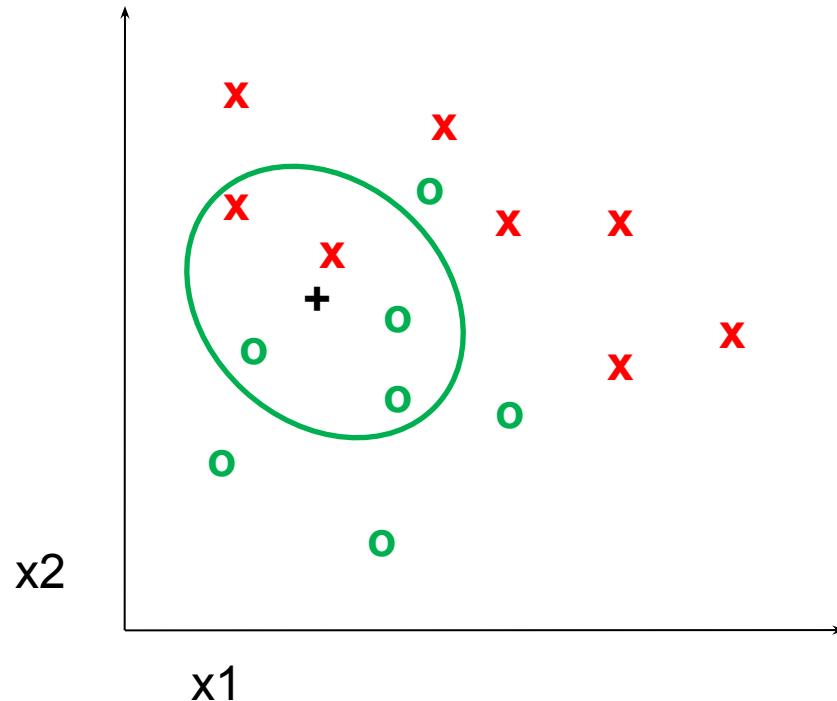
5-nearest neighbor

Distance measure (same as the ones from segmentation)

Euclidean:

$$Dist(X^n, X^m) = \sqrt{\sum_{i=1}^D (X_i^n - X_i^m)^2}$$

Where X^n and X^m are the n-th and m-th data points



Choosing the right features is important but dataset-dependent



| | Color | $D_x D_y$ | Mag-Lap | PCA Masks | PCA Gray | Cont. Greedy | Cont. DynProg | Avg. |
|--------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|--------|
| apple | 57.56% | 85.37% | 80.24% | 78.78% | 88.29% | 77.07% | 76.34% | 77.66% |
| pear | 66.10% | 90.00% | 85.37% | 99.51% | 99.76% | 90.73% | 91.71% | 89.03% |
| tomato | 98.54% | 94.63% | 97.07% | 67.80% | 76.59% | 70.73% | 70.24% | 82.23% |
| cow | 86.59% | 82.68% | 94.39% | 75.12% | 62.44% | 86.83% | 86.34% | 82.06% |
| dog | 34.63% | 62.44% | 74.39% | 72.20% | 66.34% | 81.95% | 82.93% | 67.84% |
| horse | 32.68% | 58.78% | 70.98% | 77.80% | 77.32% | 84.63% | 84.63% | 69.55% |
| cup | 79.76% | 66.10% | 77.80% | 96.10% | 96.10% | 99.76% | 99.02% | 87.81% |
| car | 62.93% | 98.29% | 77.56% | 100.0% | 97.07% | 99.51% | 100.0% | 90.77% |
| total | 64.85% | 79.79% | 82.23% | 83.41% | 82.99% | 86.40% | 86.40% | 80.87% |

Dataset: ETH-80, by B. Leibe, 2003

Results



| Category | Primary feature(s) | Secondary features |
|----------|---------------------|--------------------|
| apple | PCA Gray | Texture $D_x D_y$ |
| pear | PCA Gray / Masks | |
| tomato | Color | Texture Mag-Lap |
| cow | Texture Mag-Lap | Contour / Color |
| dog | Contour | |
| horse | Contour | |
| cup | Contour | PCA Gray / Masks |
| car | PCA Masks / Contour | Texture $D_x D_y$ |

Dataset: ETH-80, by B. Leibe, 2003

K-NN: a very useful algorithm

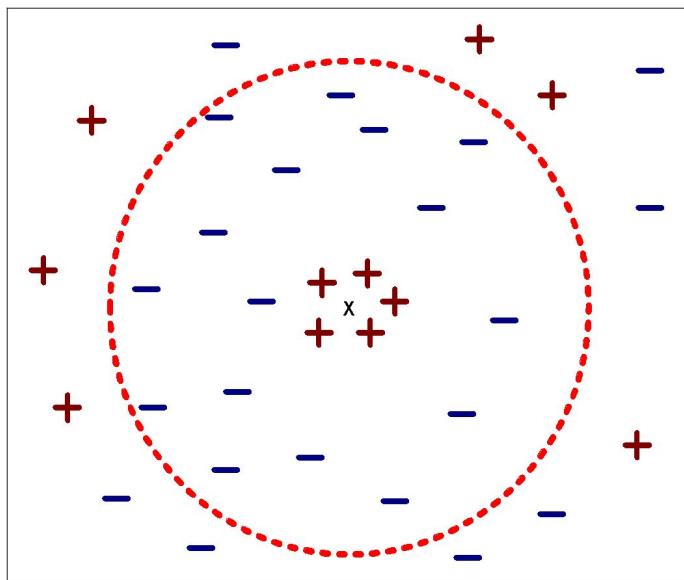
- Simple, a good one to try first
- Very flexible decision boundaries
- With infinite examples, 1-NN has a strong theoretical guarantee (out of scope for this class)

What we will learn today?

- Introduction to recognition
- A simple Object Recognition pipeline
- Choosing the right features
- A training algorithm: kNN
- Testing an algorithm
- **Challenges with kNN**
- Dimensionality reduction

K-NN: issues to keep in mind

- Choosing the value of k:
 - If too small, sensitive to noise points
 - If too large, neighborhood may include points from other classes

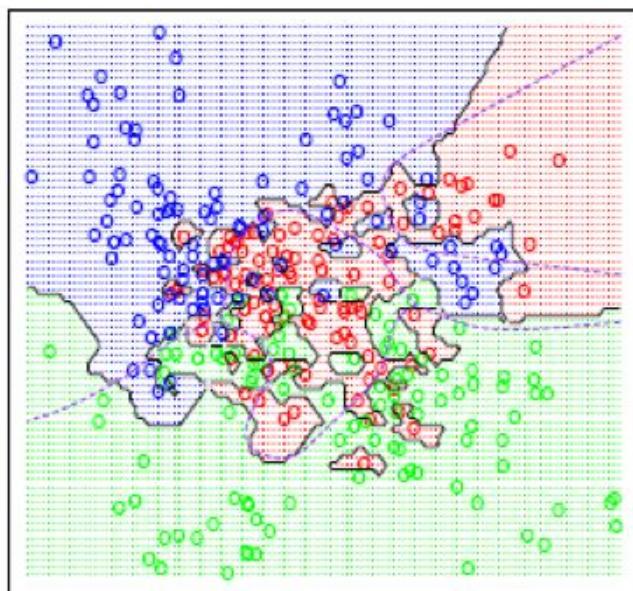


K-NN: issues to keep in mind

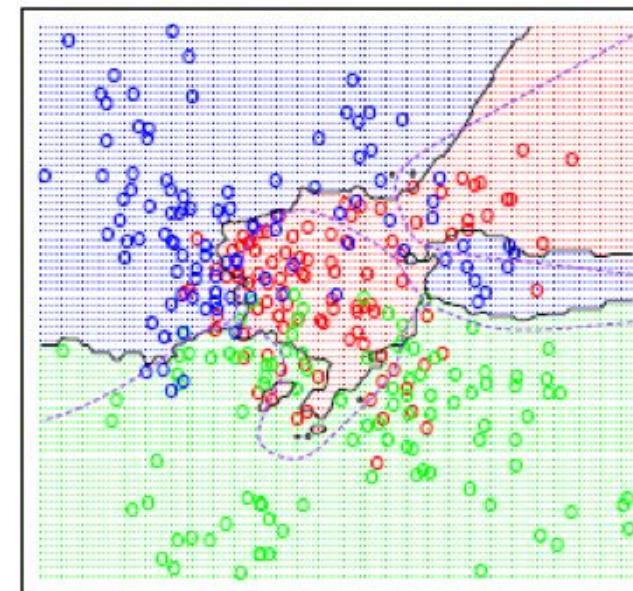
- Choosing the value of k:

- If too small, sensitive to noise points
- If too large, neighborhood may include points from other classes

K=1



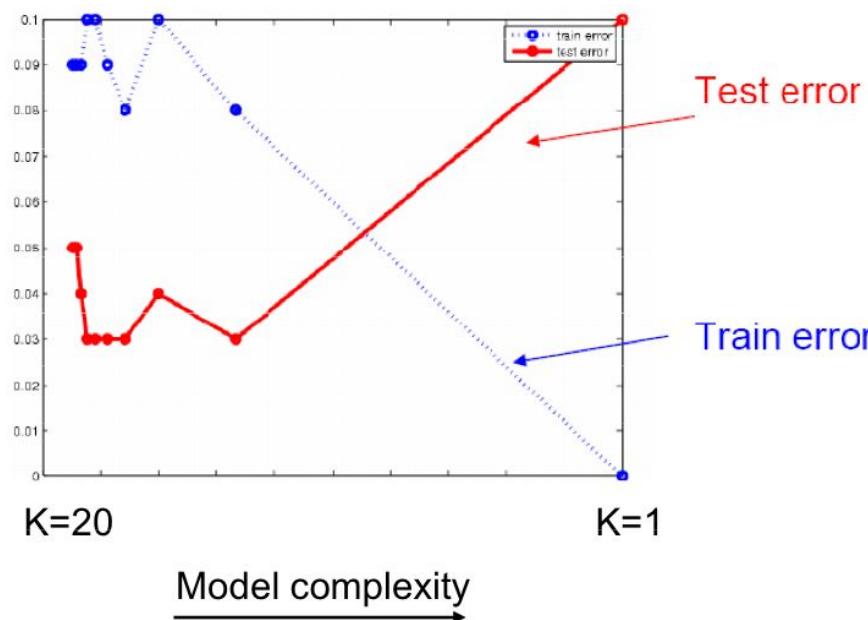
K=15



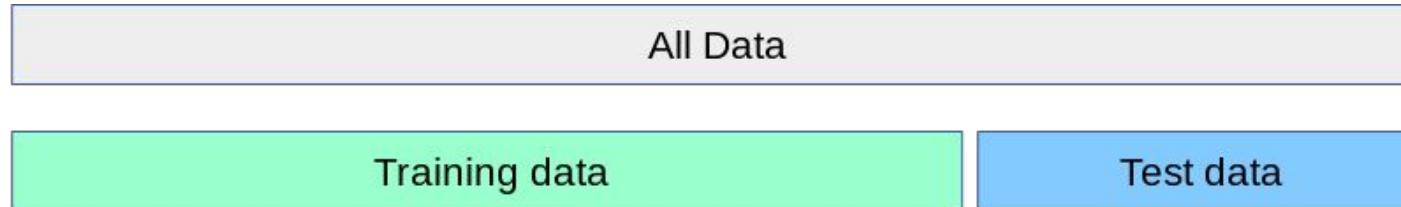
K-NN: issues to keep in mind

- Choosing the value of k:

- If too small, sensitive to noise points
- If too large, neighborhood may include points from other classes
- **Solution:** Cross validate



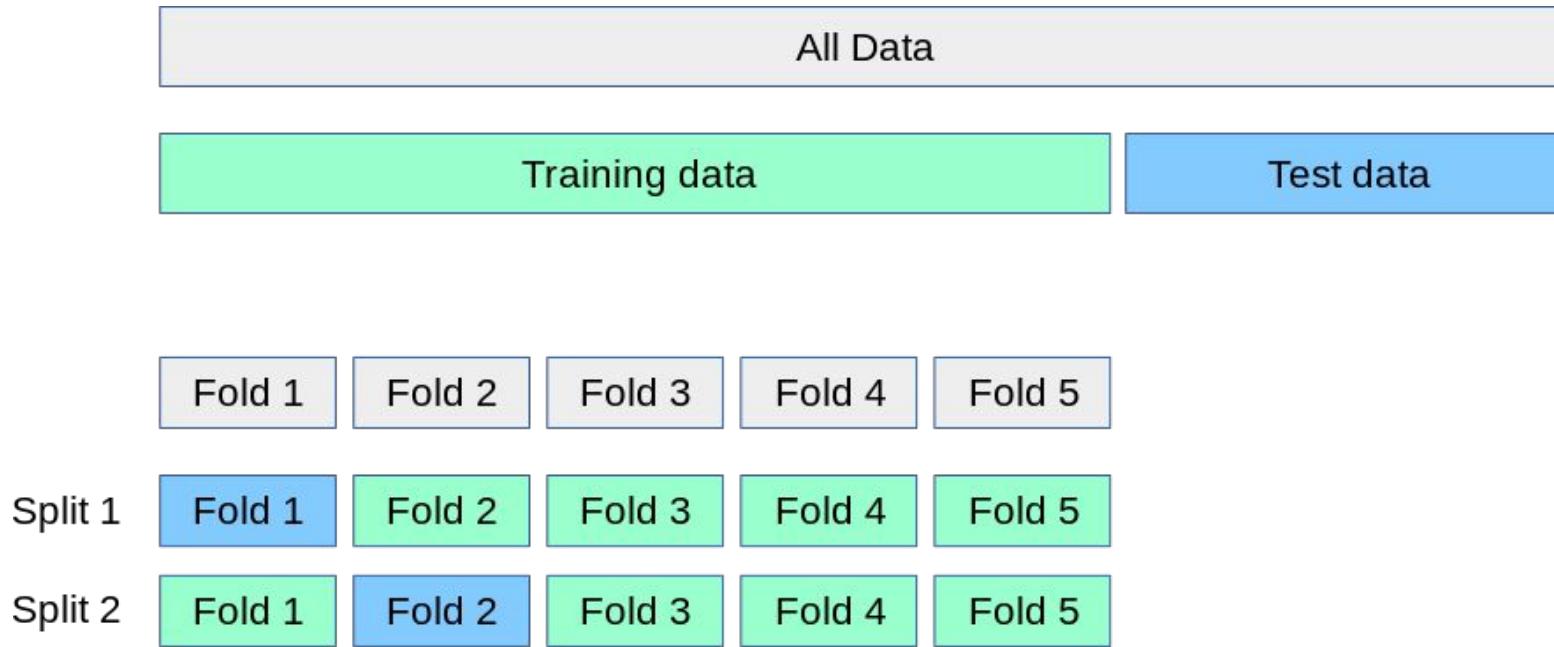
Cross validation



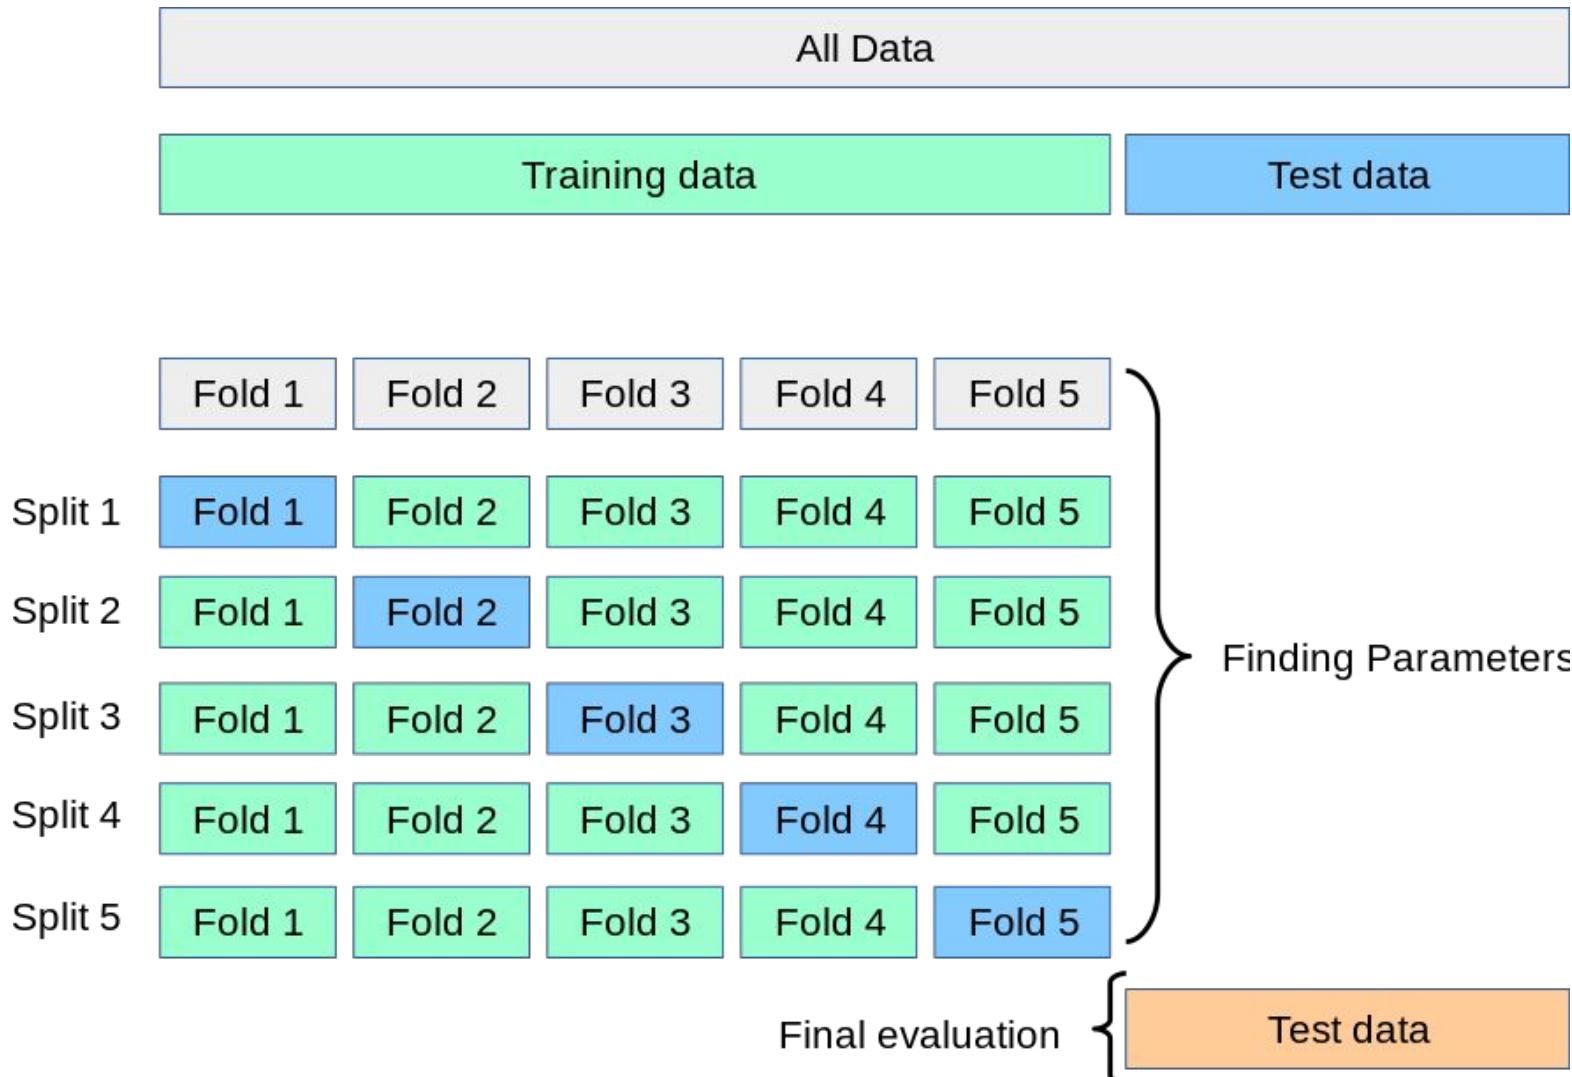
Cross validation



Cross validation



Cross validation



K-NN: issues to keep in mind

- Choosing the value of k:
 - If too small, sensitive to noise points
 - If too large, neighborhood may include points from other classes
 - **Solution:** cross validate!
- **Curse of Dimensionality**

Curse of dimensionality

- As the dimensionality increases, the number of data points required for good performance increases exponentially.
- Let's say that for a model to perform well, we need **at least 10 data points for each combination of feature values.**

Need for Data Points with Increase in Dimensions

1 Binary feature \longrightarrow 2^1 unique values \longrightarrow $2^1 \times 10 = 20$ data points

2 Binary features \longrightarrow 2^2 unique values \longrightarrow $2^2 \times 10 = 40$ data points

3 Binary features \longrightarrow 2^3 unique values \longrightarrow $2^3 \times 10 = 80$ data points

.

.

.

k Binary features \longrightarrow 2^k unique values \longrightarrow $2^k \times 10$ data points

K-NN: issues to keep in mind

- Choosing the value of k:
 - If too small, sensitive to noise points
 - If too large, neighborhood may include points from other classes
 - **Solution:** cross validate!
- Curse of Dimensionality
 - **Solution:** dimensionality reduction

What we will learn today

- Introduction to recognition
- A simple Object Recognition pipeline
- Choosing the right features
- A training algorithm: kNN
- Testing an algorithm
- Challenges with kNN
- Dimensionality reduction

Singular Value Decomposition (SVD)

$$\mathbf{U}\Sigma\mathbf{V}^T = \mathbf{A}$$

- Where \mathbf{U} and \mathbf{V} are rotation matrices, and Σ is a scaling matrix. For example:

$$U \begin{bmatrix} -.40 & .916 \\ .916 & .40 \end{bmatrix} \times \Sigma \begin{bmatrix} 5.39 & 0 \\ 0 & 3.154 \end{bmatrix} \times V^T \begin{bmatrix} -.05 & .999 \\ .999 & .05 \end{bmatrix} = A \begin{bmatrix} 3 & -2 \\ 1 & 5 \end{bmatrix}$$

Singular Value Decomposition (SVD)

- Beyond 2x2 matrices:
 - In general, if \mathbf{A} is $m \times n$, then \mathbf{U} will be $m \times m$, Σ will be $m \times n$, and \mathbf{V}^T will be $n \times n$.
 - (Note the dimensions work out to produce $m \times n$ after multiplication)

$$\begin{bmatrix} U \\ -.39 & -.92 \\ -.92 & .39 \end{bmatrix} \times \begin{bmatrix} \Sigma \\ 9.51 & 0 & 0 \\ 0 & .77 & 0 \end{bmatrix} \times \begin{bmatrix} V^T \\ -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} = \begin{bmatrix} A \\ 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

Singular Value Decomposition (SVD)

- **U** and **V** are always rotation matrices.
 - Geometric rotation may not be an applicable concept, depending on the matrix. So we call them “unitary” matrices – each column is a unit vector.
- **Σ** is a diagonal matrix
 - The number of nonzero entries = rank of **A**
 - The algorithm always sorts the entries high to low

$$\begin{bmatrix} U & \Sigma & V^T \\ \begin{bmatrix} -.39 & -.92 \\ -.92 & .39 \end{bmatrix} & \begin{bmatrix} 9.51 & 0 & 0 \\ 0 & .77 & 0 \end{bmatrix} & \begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} \end{bmatrix} = \begin{bmatrix} A \\ \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \end{bmatrix}$$

SVD Applications

- We've discussed SVD in terms of geometric transformation matrices
- But SVD of an image matrix can also be very useful
- To understand this, we'll look at a less geometric interpretation of what SVD is doing

What is SVD actually doing for images?

$$\begin{matrix} U \\ \left[\begin{matrix} -.39 & -.92 \\ -.92 & .39 \end{matrix} \right] \end{matrix} \times \begin{matrix} \Sigma \\ \left[\begin{matrix} 9.51 & 0 & 0 \\ 0 & .77 & 0 \end{matrix} \right] \end{matrix} \times \begin{matrix} V^T \\ \left[\begin{matrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{matrix} \right] \end{matrix} = \begin{matrix} A \\ \left[\begin{matrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{matrix} \right] \end{matrix}$$

- Look at how the multiplication works out, left to right:
- Column 1 of \mathbf{U} gets scaled by the first value from Σ .

$$U\Sigma = \left[\begin{matrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{matrix} \right]$$

What is SVD actually doing for images?

$$\begin{matrix} U \\ \left[\begin{matrix} -.39 & -.92 \\ -.92 & .39 \end{matrix} \right] \end{matrix} \times \begin{matrix} \Sigma \\ \left[\begin{matrix} 9.51 & 0 & 0 \\ 0 & .77 & 0 \end{matrix} \right] \end{matrix} \times \begin{matrix} V^T \\ \left[\begin{matrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{matrix} \right] \end{matrix} = \begin{matrix} A \\ \left[\begin{matrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{matrix} \right] \end{matrix}$$

- Look at how the multiplication works out, left to right:
- Column 1 of \mathbf{U} gets scaled by the first value from Σ .

$$\begin{matrix} U\Sigma \\ \left[\begin{matrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{matrix} \right] \end{matrix} \times \begin{matrix} V^T \\ \left[\begin{matrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{matrix} \right] \end{matrix}$$

What is SVD actually doing for images?

$$U \begin{bmatrix} -3.67 \\ -8.8 \end{bmatrix} \times \begin{bmatrix} 9.51 & 0 & 0 \\ 0 & .77 & 0 \end{bmatrix} \times \begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

- Look at how the multiplication works out, left to right:
- Column 1 of \mathbf{U} gets scaled by the first value from Σ .

$$U\Sigma \times \begin{bmatrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} \rightarrow A_{partial} \begin{bmatrix} 1.6 & 2.1 & 2.6 \\ 3.8 & 5.0 & 6.2 \end{bmatrix}$$

- The resulting vector gets scaled by row 1 of \mathbf{V}^T to produce a contribution to the columns of \mathbf{A}

SVD is a type dimensionality reduction

$$\begin{aligned} & \left[\begin{matrix} -3.67 \\ -8.8 \end{matrix} \right] \times \left[\begin{matrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{matrix} \right] \times \left[\begin{matrix} A_{partial} \\ 1.6 \\ 3.8 \\ 5.0 \\ 6.2 \end{matrix} \right] \\ + & \left[\begin{matrix} -3.67 \\ -8.8 \end{matrix} \right] \times \left[\begin{matrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{matrix} \right] \times \left[\begin{matrix} A_{partial} \\ -.6 \\ .2 \\ -.1 \\ 0 \\ .4 \\ -.2 \end{matrix} \right] \\ = & \left[\begin{matrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{matrix} \right] \end{aligned}$$

- Each product of (*column i of \mathbf{U}*)·(*value i from Σ*)·(*row i of \mathbf{V}^T*) produces a component of the final \mathbf{A} .

SVD is a type dimensionality reduction

$$\begin{array}{c} U\Sigma \quad V^T \\ \left[\begin{matrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{matrix} \right] \times \left[\begin{matrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{matrix} \right] \quad A_{partial} \\ \left[\begin{matrix} 1.6 & 2.1 & 2.6 \\ 3.8 & 5.0 & 6.2 \end{matrix} \right] \quad A \\ \left[\begin{matrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{matrix} \right] \end{array}$$

$$\begin{array}{c} U\Sigma \quad V^T \\ \left[\begin{matrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{matrix} \right] \times \left[\begin{matrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{matrix} \right] \quad A_{partial} \\ \left[\begin{matrix} -.6 & -.1 & .4 \\ .2 & 0 & -.2 \end{matrix} \right] \end{array}$$

- We're building \mathbf{A} as a linear combination of the columns of \mathbf{U}
- Using all columns of \mathbf{U} , we'll rebuild the original matrix perfectly
- But, in real-world data, often we can just use the first few columns of \mathbf{U} and we'll get something close (e.g. the first $\mathbf{A}_{partial}$, above)

SVD is a type dimensionality reduction

$$\begin{array}{c} U\Sigma \quad V^T \\ \left[\begin{matrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{matrix} \right] \times \left[\begin{matrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{matrix} \right] \quad A_{partial} \\ \left[\begin{matrix} 1.6 & 2.1 & 2.6 \\ 3.8 & 5.0 & 6.2 \end{matrix} \right] \quad A \\ \left[\begin{matrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{matrix} \right] \end{array}$$

$$\begin{array}{c} U\Sigma \quad V^T \\ \left[\begin{matrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{matrix} \right] \times \left[\begin{matrix} -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{matrix} \right] \quad A_{partial} \\ \left[\begin{matrix} -.6 & -.1 & .4 \\ .2 & 0 & -.2 \end{matrix} \right] \end{array}$$

- We can call those first few columns of U the **Principal Components** of the data
- They show the major patterns that can be added to produce the columns of the original matrix
- The rows of V^T show how the *principal components* are mixed to produce the columns of the matrix

SVD is a type dimensionality reduction

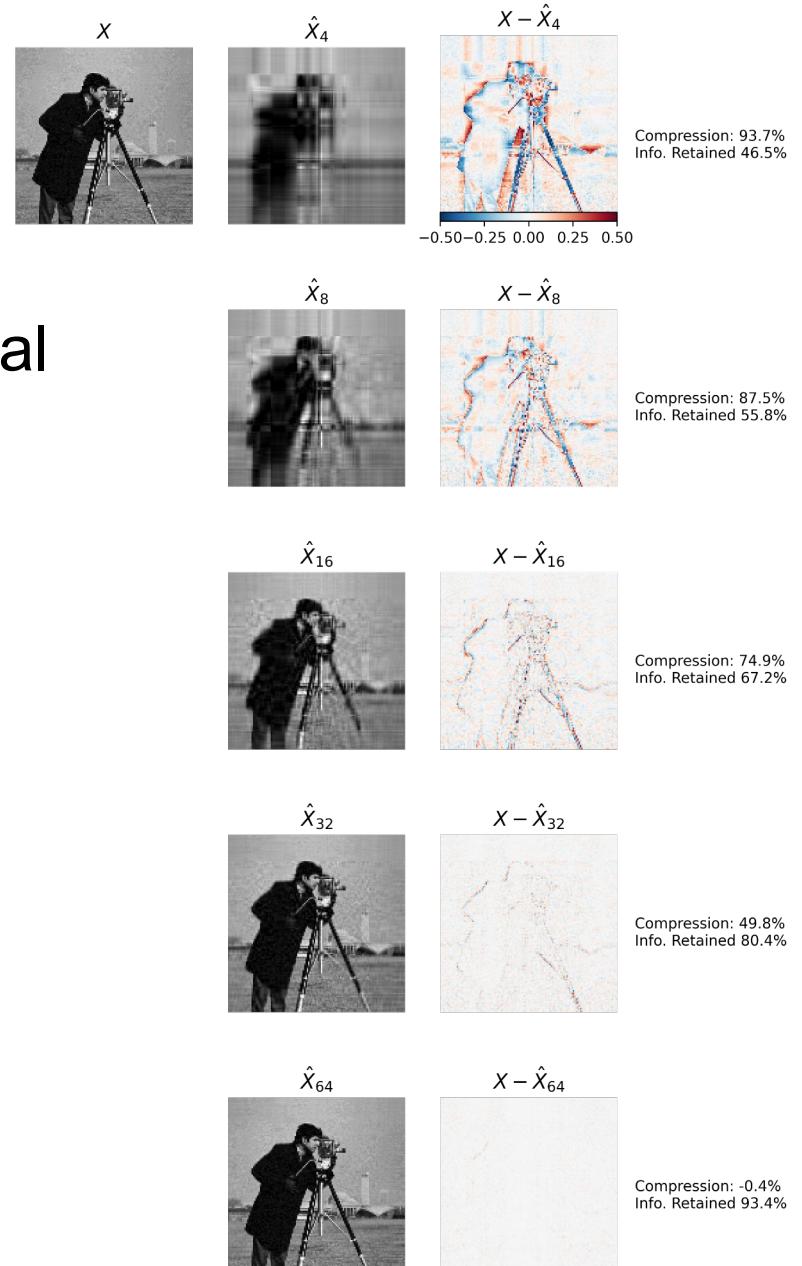
$$\begin{bmatrix} U \\ - .39 & -.92 \\ -.92 & .39 \end{bmatrix} \times \begin{bmatrix} \Sigma & & \\ 9.51 & 0 & 0 \\ 0 & 0 & .77 \end{bmatrix} \times \begin{bmatrix} V^T \\ -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} = \begin{bmatrix} A \\ 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

We can look at Σ to see that the first column has a large effect

while the second column has a much smaller effect in this example

Image compression

- For this image, using **only the first 16** of 300 principal components produces a recognizable reconstruction
- Using the first 64 almost perfectly reconstructs the image



SVD for symmetric matrices

- If A is a symmetric matrix, it can be decomposed as the following:
- Compared to a traditional $A = \Phi\Sigma\Phi^T$ Φ is an orthogonal matrix.

What we have learned today?

- Introduction to recognition
- A simple Object Recognition pipeline
- Choosing the right features
- A training algorithm: kNN
- Testing an algorithm
- Challenges with kNN
- Dimensionality reduction

Next lecture

PCA + LDA

Extra slides (out of scope)

for those of you curious about how SVD is calculated
and what PCA is usually used for outside of computer vision

Principal Component Analysis

$$\begin{bmatrix} -3.67 & U\Sigma & V^T \\ -8.8 & -.71 & -.42 & -.57 & -.70 \\ & .30 & .81 & .11 & -.58 \\ & 0 & .41 & -.82 & .41 \end{bmatrix} \times \begin{bmatrix} A_{partial} \\ 1.6 \\ 3.8 \\ 2.1 \\ 5.0 \\ 2.6 \\ 6.2 \end{bmatrix}$$

- Remember, columns of \mathbf{U} are the *Principal Components* of the data: the major patterns that can be added to produce the columns of the original matrix
- One use of this is to construct a matrix where each column is a separate data sample
- Run SVD on that matrix, and look at the first few columns of \mathbf{U} to see patterns that are common among the columns
- This is called *Principal Component Analysis* (or PCA) of the data samples

Principal Component Analysis

$$\begin{bmatrix} -3.67 & -.71 & 0 \\ -8.8 & .30 & 0 \end{bmatrix} \times \begin{bmatrix} V^T \\ -.42 & -.57 & -.70 \\ .81 & .11 & -.58 \\ .41 & -.82 & .41 \end{bmatrix} = A_{partial} \begin{bmatrix} 1.6 & 2.1 & 2.6 \\ 3.8 & 5.0 & 6.2 \end{bmatrix}$$

- Often, raw data samples have a lot of redundancy and patterns
- PCA can allow you to represent data samples as weights on the principal components, rather than using the original raw form of the data
- By representing each sample as just those weights, you can represent just the “meat” of what’s different between samples.
- This minimal representation makes machine learning and other algorithms much more efficient

How is SVD computed?

- For this class: tell PYTHON to do it. Use the result.
- But, if you're interested, one computer algorithm to do it makes use of Eigenvectors!

Eigenvector definition

- Suppose we have a square matrix \mathbf{A} . We can solve for vector x and scalar λ such that $\mathbf{Ax} = \lambda x$
- In other words, find vectors where, if we transform them with \mathbf{A} , the only effect is to scale them with no change in direction.
- These vectors are called eigenvectors (German for “self vector” of the matrix), and the scaling factors λ are called eigenvalues
- An $m \times m$ matrix will have $\leq m$ eigenvectors where λ is nonzero

Finding eigenvectors

- Computers can find an x such that $Ax = \lambda x$ using this iterative algorithm:
 - $X = \text{random unit vector}$
 - while(x hasn't converged)
 - $X = Ax$
 - normalize x
- x will quickly converge to an eigenvector
- Some simple modifications will let this algorithm find all eigenvectors

Finding SVD

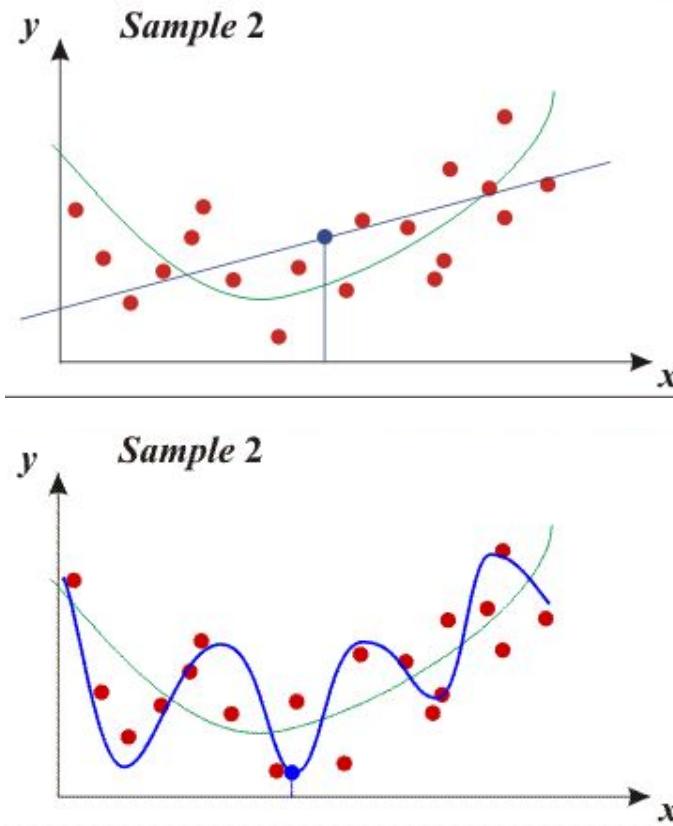
- Eigenvectors are for square matrices, but SVD is for all matrices
- To do $\text{svd}(A)$, computers can do this:
 - Take eigenvectors of AA^T (matrix is always square).
 - These eigenvectors are the columns of \mathbf{U} .
 - Square root of eigenvalues are the singular values (the entries of Σ).
 - Take eigenvectors of A^TA (matrix is always square).
 - These eigenvectors are columns of \mathbf{V} (or rows of \mathbf{V}^T)

Finding SVD

- Moral of the story: SVD is fast, even for large matrices
- It's useful for a lot of stuff
- There are also other algorithms to compute SVD or part of the SVD
 - Python's `np.linalg.svd()` command has options to efficiently compute only what you need, if performance becomes an issue

A detailed geometric explanation of SVD is here:
<http://www.ams.org/samplings/feature-column/fcarc-svd>

Bias-Variance Trade-off

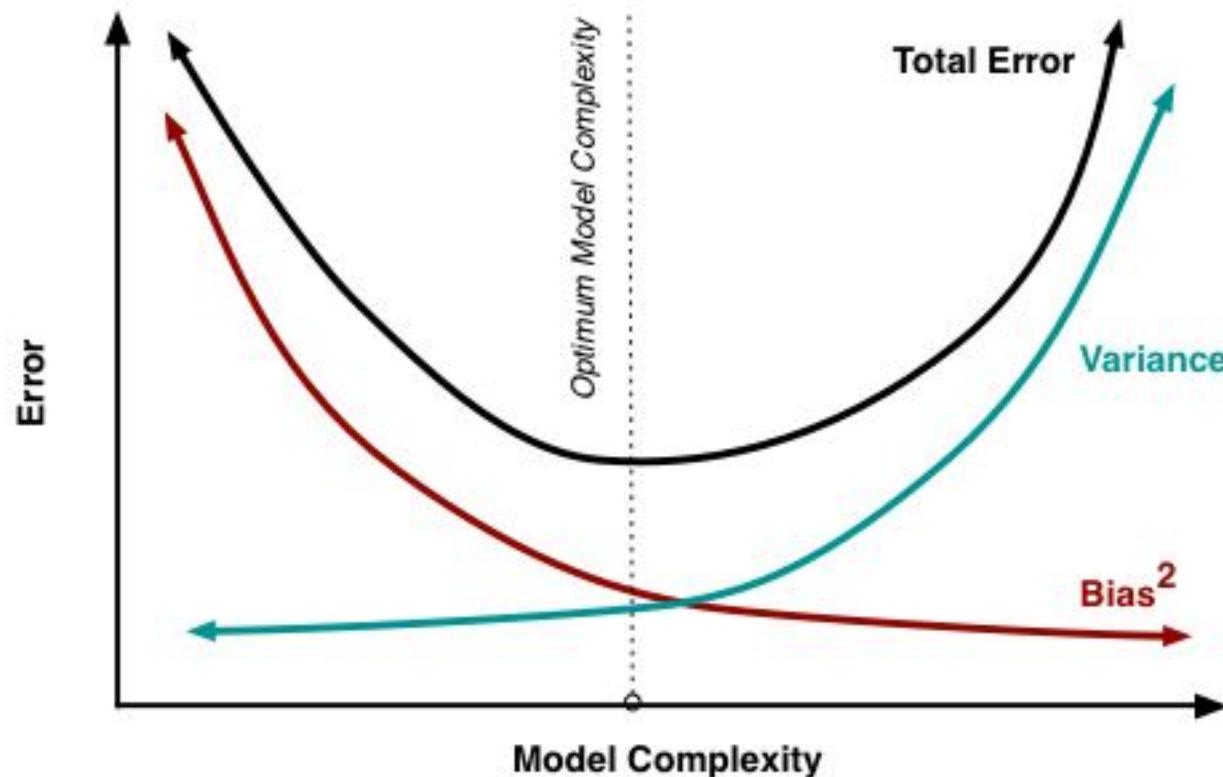


- Models with too few parameters are inaccurate because of a large bias (not enough flexibility).
- Models with too many parameters are inaccurate because of a large variance (too much sensitivity to the sample).

Bias versus variance

- Components of generalization error
 - **Bias**: how much the average model over all training sets differ from the true model?
 - Error due to inaccurate assumptions/simplifications made by the model
 - **Variance**: how much models estimated from different training sets differ from each other
- **Underfitting**: model is too “simple” to represent all the relevant class characteristics
 - High bias and low variance
 - High training error and high test error
- **Overfitting**: model is too “complex” and fits irrelevant characteristics (noise) in the data
 - Low bias and high variance
 - Low training error and high test error

Bias versus variance trade off



No Free Lunch Theorem



In a supervised learning setting, we can't tell which classifier will have best generalization

Remember...

- No classifier is inherently better than any other: you need to make assumptions to generalize
- Three kinds of error
 - Inherent: unavoidable
 - Bias: due to over-simplifications
 - Variance: due to inability to perfectly estimate parameters from limited data



How to reduce variance?

- Choose a simpler classifier
- Regularize the parameters
- Get more training data

How do you reduce bias?

Last remarks about applying machine learning methods to object recognition

- There are machine learning algorithms to choose from
- Know your data:
 - How much supervision do you have?
 - How many training examples can you afford?
 - How noisy?
- Know your goal (i.e. task):
 - Affects your choices of representation
 - Affects your choices of learning algorithms
 - Affects your choices of evaluation metrics
- Understand the math behind each machine learning algorithm under consideration!

PCA by SVD

- An alternative manner to compute the principal components, based on singular value decomposition
- Quick reminder: SVD
 - Any real $n \times m$ matrix ($n > m$) can be decomposed as

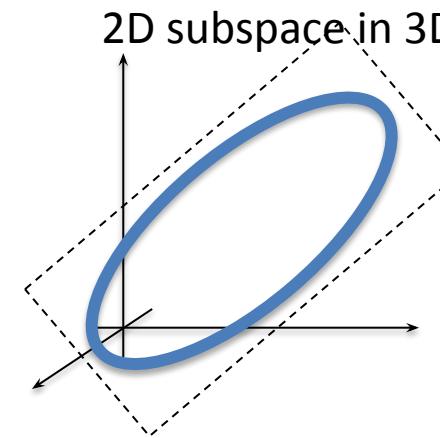
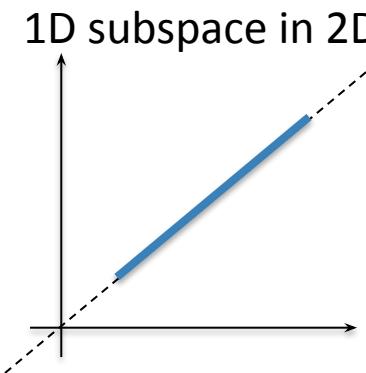
$$A = M \Pi N^T$$

- Where M is an $(n \times m)$ column orthonormal matrix of left singular vectors (columns of M)
- Π is an $(m \times m)$ diagonal matrix of singular values
- N^T is an $(m \times m)$ row orthonormal matrix of right singular vectors (columns of N)

$$M^T M = I \quad N^T N = I$$

PCA Formulation

- Basic idea:
 - If the images (or their features) live in a subspace, it is going to look very flat when viewed from the full feature space, e.g.



Slide inspired by N. Vasconcelos

Alternative PCA Formulation

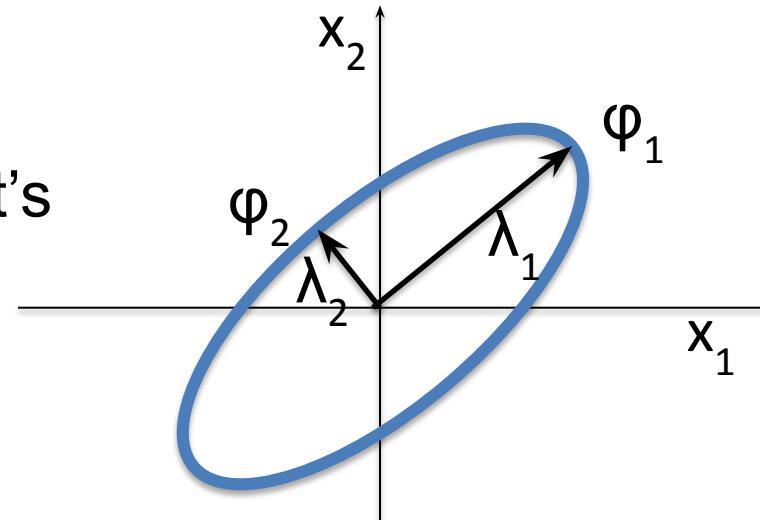
- Assume images \mathbf{x} is Gaussian with covariance Σ .
- Recall that a gaussian is defined with it's mean and variance:

$$\mathbf{X} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$$

- Recall that $\boldsymbol{\mu}$ and $\boldsymbol{\Sigma}$ of a gaussian are defined as:

$$\boldsymbol{\mu} = \mathbb{E}[\mathbf{X}]$$

$$\boldsymbol{\Sigma} =: \mathbb{E}[(\mathbf{X} - \boldsymbol{\mu})(\mathbf{X} - \boldsymbol{\mu})^T] = [\text{Cov}[X_i, X_j]; 1 \leq i, j \leq k]$$



Alternative PCA formulation

- Since gaussians are symmetric, it's covariance matrix is also a symmetric matrix. So we can express it as:
 - $\Sigma = \mathbf{U}\Lambda\mathbf{U}^T = \mathbf{U}\Lambda^{1/2}(\mathbf{U}\Lambda^{1/2})^T$

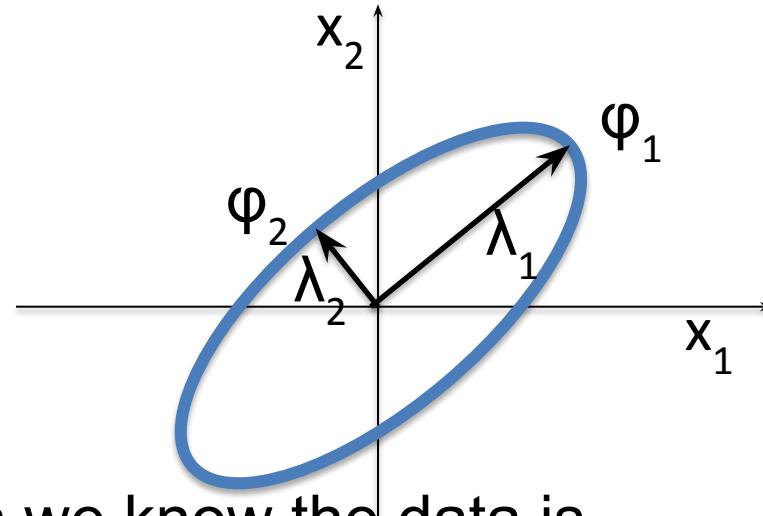
$$\mathbf{X} \sim \mathcal{N}(\boldsymbol{\mu}, \Sigma) \iff \mathbf{X} \sim \boldsymbol{\mu} + \mathbf{U}\Lambda^{1/2}\mathcal{N}(0, \mathbf{I})$$

$$\iff \mathbf{X} \sim \boldsymbol{\mu} + \mathbf{U}\mathcal{N}(0, \boldsymbol{\Lambda}).$$

Alternative PCA Formulation

- If x is Gaussian with covariance Σ ,

- Principal components ϕ_i are the eigenvectors of Σ
- Principal lengths λ_i are the eigenvalues of Σ



- by computing the eigenvalues we know the data is
 - Not flat if $\lambda_1 \approx \lambda_2$
 - Flat if $\lambda_1 \gg \lambda_2$

Alternative PCA Algorithm (training)

► Given sample $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$, $\mathbf{x}_i \in \mathcal{R}^d$

- compute sample mean: $\hat{\mu} = \frac{1}{n} \sum_i (\mathbf{x}_i)$
- compute sample covariance: $\hat{\Sigma} = \frac{1}{n} \sum_i (\mathbf{x}_i - \hat{\mu})(\mathbf{x}_i - \hat{\mu})^T$

- compute eigenvalues and eigenvectors of $\hat{\Sigma}$

$$\hat{\Sigma} = \Phi \Lambda \Phi^T, \quad \Lambda = \text{diag}(\sigma_1^2, \dots, \sigma_n^2) \quad \Phi^T \Phi = I$$

- order eigenvalues $\sigma_1^2 > \dots > \sigma_n^2$

- if, for a certain k , $\sigma_k \ll \sigma_1$ eliminate the eigenvalues and eigenvectors above k .

Alternative PCA Algorithm (testing)

- ▶ Given principal components $\phi_i, i \in 1, \dots, k$ and a test sample $\mathcal{T} = \{t_1, \dots, t_n\}, t_i \in \mathcal{R}^d$

- subtract mean to each point $t'_i = t_i - \hat{\mu}$

- project onto eigenvector space $y_i = At'_i$ where

$$A = \begin{bmatrix} \phi_1^T \\ \vdots \\ \phi_k^T \end{bmatrix}$$

- use $\mathcal{T}' = \{y_1, \dots, y_n\}$ to estimate class conditional densities and do all further processing on y .